

# Multivariate Linear Regression

## General Principles

To study relationships between multiple continuous independent variables (e.g., the effect of weight and age on height), we can use a multiple regression approach. Essentially, we extend [Linear Regression for continuous variable](#) by adding a regression coefficient  $\beta_x$  for each continuous variable (e.g.,  $\beta_{weight}$  and  $\beta_{age}$ ).

## Considerations

### Note

- We have the same considerations as for the [Regression for continuous variable](#).
- The model interpretation of the regression coefficients  $\beta_x$  is considered for fixed values of the other independent variable(s)' regression coefficients—i.e., for a given age,  $\beta_{weight}$  represents the expected change in the dependent variable (height) for each one-unit increase in weight, holding all other variables (e.g., age) constant.

## Example

Below is example code demonstrating Bayesian multiple linear regression using the Bayesian Inference (BI) package. Data consist of three continuous variables (*height*, *weight*, *age*), and the goal is to estimate the effect of *weight* and *age* on *height*. This example is based on McElreath (2018).

## Python

```

from BI import bi

# Setup device-----
m = bi(platform='cpu')

# Import Data & Data Manipulation -----
from importlib.resources import files
# Import
data_path = m.load.howell1(only_path = True)
m.data(data_path, sep=';')
m.df = m.df[m.df.age > 18] # Subset data to adults
m.scale(['weight', 'age']) # Normalize

# Define model -----
def model(height, weight, age):
    # Parameter prior distributions
    alpha = m.dist.normal(0, 0.5, name = 'alpha')
    beta1 = m.dist.normal(0, 0.5, name = 'beta1')
    beta2 = m.dist.normal(0, 0.5, name = 'beta2')
    sigma = m.dist.uniform(0, 50, name = 'sigma')
    # Likelihood
    m.dist.normal(alpha + beta1 * weight + beta2 * age, sigma, obs = height)

# Run MCMC -----
m.fit(model) # Optimize model parameters through MCMC sampling

# Summary -----
m.summary()

```

jax.local\_device\_count 32

0%| 0/1000 [00:00<?, ?it/s] warmup: 0%| 1/1000 [00:00<09:45, 1.71it/s]  
arviz - WARNING - Shape validation failed: input\_shape: (1, 500), minimum\_shape: (chains=2, 500)

	mean	sd	hdi_5.5%	hdi_94.5%	mcse_mean	mcse_sd	ess_bulk	ess_tail	r_hat
alpha	5.20	0.49	4.46	6.06	0.02	0.02	469.94	349.45	NaN
beta1	0.20	0.51	-0.60	1.03	0.02	0.03	570.81	264.55	NaN
beta2	-0.02	0.49	-0.89	0.69	0.02	0.02	576.00	338.15	NaN
sigma	49.98	0.02	49.96	50.00	0.00	0.00	579.62	271.13	NaN

## R

```
library(BayesianInference)
m=importBI(platform='cpu')

# Import Data & Data Manipulation -----
m$data(m$load$howell1(only_path = T), sep=';')# Import
m$df = m$df[m$df$age > 18,] # Subset data to adults
m$scale(list('weight', 'age')) # Normalize
m$data_to_model(list('weight', 'height', 'age')) # Send to model (convert to jax array)

# Define model -----
model <- function(height, weight, age){
  # Parameter prior distributions
  alpha = bi.dist.normal(0, 0.5, name = 'a')
  beta1 = bi.dist.normal(0, 0.5, name = 'b1')
  beta2 = bi.dist.normal(0, 0.5, name = 'b2')
  sigma = bi.dist.uniform(0, 50, name = 's')
  # Likelihood
  bi.dist.normal(alpha + beta1 * weight + beta2 * age, sigma, obs=height)
}

# Run MCMC -----
m$fit(model) # Optimize model parameters through MCMC sampling

# Summary -----
m$summary() # Get posterior distributions
```

### 🔥 Caution

For R users, if you create the regression coefficient in a single call:

```
betas = bi.dist.normal(0, 0.5, name = 'regression_coefficients', shape = (2,))
```

you need to index them starting by 0:

```
m$normal(alpha + betas[0] * weight + betas[1] * age, sigma, obs=height)
```

## Mathematical Details

### Frequentist formulation

We model the relationship between the independent variables  $(X_{1i}, X_{2i}, \dots, X_{[K,i]})$  and the dependent variable  $Y$  using the following equation:

$$Y_i = \alpha + \beta_1 X_{[1,i]} + \beta_2 X_{[2,i]} + \dots + \beta_n X_{[K,i]} + \epsilon_i$$

Where:

- $Y_i$  is the dependent variable for observation  $i$ .
- $\alpha$  is the intercept term.
- $X_{[1,i]}, X_{[2,i]}, \dots, X_{[K,i]}$  are the values of the independent variables for observation  $i$ .
- $\beta_1, \beta_2, \dots, \beta_K$  are the regression coefficients.
- $\epsilon_i$  is the error term for observation  $i$ , and the vector of the error terms,  $\epsilon$ , are assumed to be independent and identically distributed.

### Bayesian formulation

In the Bayesian formulation, we define each parameter with priors . We can express the Bayesian model as follows:

$$Y_i \sim \text{Normal}(\alpha + \sum_{k=1}^K \beta_k X_{[K,i]}, \sigma)$$

$$\alpha \sim \text{Normal}(0, 1)$$

$$\beta_k \sim \text{Normal}(0, 1)$$

$$\sigma \sim \text{Uniform}(0, 50)$$

Where:

- $Y_i$  is the dependent variable for observation  $i$ .
- $\alpha$  is the intercept term, which in this case has a unit-normal prior.

- $\beta_k$  are slope coefficients for the  $K$  distinct independent variables, which also have unit-normal priors.
- $X_{[1,i]}, X_{[2,i]}, \dots, X_{[K,i]}$  are the values of the independent variables for observation  $i$ .
- $\sigma$  is a standard deviation parameter, which here has a Uniform prior that constrains it to be positive.

## Reference(s)

McElreath, Richard. 2018. *Statistical Rethinking: A Bayesian course with examples in R and Stan*. Chapman; Hall/CRC.