

Varying Intercepts Models

General Principles

To model the relationship between a dependent variable and an independent variable while allowing for different intercepts across groups or clusters, we can use a *Varying Intercepts* model. This approach is particularly useful when data is grouped (e.g., by subject, location, or time period) and we expect the baseline level of the outcome to vary across these groups.

Considerations

🔥 Caution

- We have the same considerations as for [Regression for a continuous variable](#).
- The main idea of varying intercepts is to generate an intercept for each group, allowing each group to start at different levels. Thus, the intercept α_k is defined based on the k declared groups.
- Each intercept has its own prior - i.e., a hyper-prior .
- In the code below, the *hyper-prior* is `a_bar`.

Example

Below is an example code snippet demonstrating Bayesian regression with varying intercepts using the Bayesian Inference (BI) package. The data consists of a dependent variable representing individuals' survival (`surv`) and an independent categorical variable (`tank`), which indicates the tank where the individual was born, with a total of 48 tanks. This example is based on McElreath (2018).

Python

```
from BI import bi
import numpy as np

# Setup device-----
m = bi(platform='cpu')

# Import Data & Data Manipulation -----
# Import
from importlib.resources import files
data_path = files('BI.resources.data') / 'reedfrogs.csv'
m.data(data_path, sep=';')
# Manipulate
m.df["tank"] = np.arange(m.df.shape[0])

# Define model -----
def model(tank, surv, density):
    sigma = m.dist.exponential( 1, name = 'sigma')
    a_bar = m.dist.normal( 0., 1.5, name = 'a_bar')
    alpha = m.dist.normal( a_bar, sigma, shape= tank.shape, name = 'alpha')
    p = alpha[tank]
    m.dist.binomial(total_count = density, logits = p, obs=surv)

# Run sampler -----
m.fit(model)

# Diagnostic -----
m.summary()
```

R

```
library(BI)

# setup platform-----
m=importbi(platform='cpu')

# Import data -----
m$data=paste(system.file(package = "BI"),"/data/reedfrogs.csv", sep = ''), sep=';')
m$df$tank = c(0:(nrow(m$df)-1)) # Manipulate
m$data_to_model(list('tank', 'surv', 'density')) # Manipulate
```

```

m$data_on_model$tank = m$data_on_model$tank$astype(jnp$int32) # Manipulate
m$data_on_model$surv = m$data_on_model$surv$astype(jnp$int32) # Manipulate

# Define model -----
model <- function(tank, surv, density){
  # Parameter prior distributions
  sigma = bi.dist.exponential( 1, name = 'sigma', shape=c(1))
  a_bar = bi.dist.normal(0, 1.5, name='a_bar', shape=c(1))
  alpha = bi.dist.normal(a_bar, sigma, name='alpha', shape =c(48))
  p = alpha[tank]
  # Likelihood
  m$binomial(total_count = density, logits = p, obs=surv)
}

# Run MCMC -----
m$run(model) # Optimize model parameters through MCMC sampling

# Summary -----
m$summary() # Get posterior distribution

```

Mathematical Details

Frequentist formulation

We model the relationship between the independent variable X and the outcome variable Y with varying intercepts α for each group k using the following equation:

$$Y_{ik} = \alpha_k + \beta X_{ik} + \sigma$$

Where:

- Y_{ik} is the outcome variable for observation i in group k .
- α_k is the varying intercept for group k .
- X_{ik} is the independent variable for observation i in group k .
- β is the regression coefficient.
- σ is the error term, typically assumed to be normally distributed and positive.

Bayesian Model

We can express the Bayesian regression model accounting for priors distributions as follows:

$$\begin{aligned}Y_{ik} &\sim \text{Normal}(\mu_{ik}, \sigma) \\ \mu_{ik} &= \alpha_k + \beta X_{ik} \\ \alpha_k &\sim \text{Normal}(\mu_{\alpha_k}, \sigma_{\alpha_k}) \\ \beta &\sim \text{Normal}(0, 1) \\ \sigma &\sim \text{Exponential}(1) \\ \mu_{\alpha_k} &\sim \text{Normal}(0, 1) \\ \sigma_{\alpha_k} &\sim \text{Exponential}(1)\end{aligned}$$

Where:

- Y_{ik} is the outcome variable for observation i in group k .
- α_k is the varying intercept for group k .
- μ_{α_k} is the overall mean intercept.
- σ_{α_k} is the variance of the intercepts across groups.
- β is the regression coefficient.
- σ is the standard deviation of the error term.

Notes

Note

- We can apply multiple variables similarly to [Chapter 2](#).
- We can apply interaction terms similarly to [Chapter 3](#).
- We can apply categorical variables similarly to [Chapter 4](#).
- We can apply varying intercepts with any distribution developed in previous chapters.

Reference(s)

McElreath, Richard. 2018. *Statistical Rethinking: A Bayesian Course with Examples in r and Stan*. Chapman; Hall/CRC.