

BNN for Classification

General Principles

To model complex, non-linear boundaries between classes, we can extend the [Bayesian Neural Network \(BNN\)](#) architecture to classification tasks. While a regression BNN outputs a continuous value modeled with a Normal likelihood, a classification BNN outputs probabilities.

For binary classification, the final layer of the network condenses the hidden representations into a single output score. This score is then passed through a **sigmoid** activation (or used directly as the **logits** parameter) to map it to a probability between 0 and 1. We then use a **Bernoulli** or **Binomial** distribution to model the likelihood of the observed binary outcomes.

Considerations

Note

- **Uncertainty in Classification:** In standard neural networks, softmax or sigmoid outputs are often overconfident. By using a BNN, we obtain a posterior predictive distribution that reflects true uncertainty, leading to better-calibrated probabilities. Areas with little training data will appropriately show high uncertainty rather than confident misclassifications.
- **Activation and Likelihood:** For binary classification, we pair the linear output of the network's final layer with a Binomial or Bernoulli likelihood. The final linear output acts as the **logit** input to the likelihood function, effectively applying a link function to constrain outputs to $[0, 1]$.
- **Prior distributions:** Similar to regression BNNs, we apply weakly-informative priors (like $\text{Normal}(0, 1)$) to the weights of each layer.

Example

Below is an example code snippet demonstrating a *Bayesian Neural Network for classification* using the Bayesian Inference (**BI**) package. This example is inspired by stochastic variational inference tutorials and uses a synthetic nested moons dataset (typically generated via `make_moons`).

Python

```
from BI import bi
import jax.numpy as jnp
from sklearn.datasets import make_moons

# Setup device-----
m = bi(platform='cpu')

# Generate Synthetic Data -----
# Two interleaving half-moon shapes
X, Y = make_moons(n_samples=500, noise=0.25, random_state=42)

# Convert to JAX arrays
X = jnp.array(X)
Y = jnp.array(Y)

m.data_on_model = dict(X=X, Y=Y)

# Define model -----
def model(X, Y, D_H1=4, D_H2=3):
    N, D_X = X.shape

    # First hidden layer: 2 input features -> 4 hidden units
    w1 = m.bnn.layer_linear(
        X,
        dist=m.dist.normal(0, 1, name='w1', shape=(D_X, D_H1)),
        activation='tanh'
    )

    # Second hidden layer: 4 hidden units -> 3 hidden units
    w2 = m.bnn.layer_linear(
        X=w1,
        dist=m.dist.normal(0, 1, name='w2', shape=(D_H1, D_H2)),
```

```

        activation='tanh'
    )

    # Final output layer: 3 hidden units -> 1 output
    w3 = m.bnn.layer_linear(
        X=w2,
        dist=m.dist.normal(0, 1, name='w3', shape=(D_H2, 1))
    )

    # Squeeze the final output to match Y's shape of (N,)
    logits = w3.squeeze(-1)

    # Likelihood mapping the logits to binary outcomes
    m.dist.binomial(total_count=1, logits=logits, obs=Y, name='Y')

# Run mcmc -----
m.fit(model, progress_bar=False) # Approximate posterior distributions

# Predictions from the model -----
import matplotlib.pyplot as plt

# Create a grid to evaluate the model
n_grid = 50
x0 = jnp.linspace(X[:, 0].min() - 0.5, X[:, 0].max() + 0.5, n_grid)
x1 = jnp.linspace(X[:, 1].min() - 0.5, X[:, 1].max() + 0.5, n_grid)
xx0, xx1 = jnp.meshgrid(x0, x1)
X_grid = jnp.c_[xx0.ravel(), xx1.ravel()]

# Swap data on model temporarily to predict on the grid
m.data_on_model = dict(X=X_grid, Y=jnp.zeros(X_grid.shape[0]))
pred = m.sample(samples=500)['Y']
p_mean = jnp.mean(pred, axis=0)

# Plotting the posterior predictive mean
fig, ax = plt.subplots(figsize=(8, 6), constrained_layout=True)
contour = ax.contourf(xx0, xx1, p_mean.reshape(n_grid, n_grid), cmap="coolwarm", alpha=0.6)
scatter = ax.scatter(X[:, 0], X[:, 1], c=Y, cmap="coolwarm", edgecolors='k')
ax.set(title="Posterior Predictive Mean", xlabel="Feature 1", ylabel="Feature 2")
fig.colorbar(contour, ax=ax)

```

jax.local_device_count 16


This function is still in development. Use it with caution.

This function is still in development. Use it with caution.
This function is still in development. Use it with caution.
This function is still in development. Use it with caution.
This function is still in development. Use it with caution.
This function is still in development. Use it with caution.
This function is still in development. Use it with caution.
This function is still in development. Use it with caution.
This function is still in development. Use it with caution.

C:\Users\Sosa\Documents\BI\BI\Main\main.py:513: UserWarning:

Sample's batch dimension size 4000 is different from the provided 500 num_samples argument. l

This function is still in development. Use it with caution.
This function is still in development. Use it with caution.
This function is still in development. Use it with caution.



Julia

```
using BayesianInference
using PyCall

# Setup device-----
m = importBI(platform="cpu")

# Generate Synthetic Data using Python's scikit-learn -----
sk_datasets = pyimport("sklearn.datasets")
X = sk_datasets.make_moons(n_samples=500, noise=0.25, random_state=42)[1]
Y = sk_datasets.make_moons(n_samples=500, noise=0.25, random_state=42)[2]

m.data_on_model["X"] = X
m.data_on_model["Y"] = Y

# Define model -----
@BI function model(X, Y)
    N, D_X = size(X)
    D_H1 = 4
    D_H2 = 3

    # First hidden layer
    w1 = m.bnn.layer_linear(
        X,
        dist=m.dist.normal(0, 1, name="w1", shape=(D_X, D_H1)),
        activation="tanh"
    )

    # Second hidden layer
    w2 = m.bnn.layer_linear(
        w1,
        dist=m.dist.normal(0, 1, name="w2", shape=(D_H1, D_H2)),
        activation="tanh"
    )

    # Final output layer
    w3 = m.bnn.layer_linear(
        w2,
        dist=m.dist.normal(0, 1, name="w3", shape=(D_H2, 1))
    )
end
```

```

# Extract logits
logits = w3[:, 1]

# Likelihood mapping the logits to binary outcomes
m.dist.binomial(total_count=1, logits=logits, obs=Y, name="Y")
end

# Run mcmc -----
m.fit(model, num_samples=500, progress_bar=false)

```

Mathematical Details

«««< HEAD

Bayesian Formulation

For a binary classification task spanning N observations, we model the probability p_i that the response $Y_i \in \{0, 1\}$ is 1. The input features X are passed through a series of hidden layers mapping to this probability:

1. **First Hidden Layer:**

$$\Omega_{i,1} = \tanh(X_i \Theta_1)$$

Where X_i is a 2-vector of features and Θ_1 is a 2×4 weight matrix parameterized by a Standard Normal prior: $\Theta_1 \sim \text{Normal}(0, 1)$.

2. **Second Hidden Layer:**

$$\Omega_{i,2} = \tanh(\Omega_{i,1} \Theta_2)$$

Where Θ_2 is a 4×3 weight matrix parameterized by a Standard Normal prior: $\Theta_2 \sim \text{Normal}(0, 1)$.

3. **Output Layer (Logits):**

$$\text{logit}(p_i) = \Omega_{i,2} \Theta_3$$

Where Θ_3 is a 3×1 weight matrix parameterized by a Standard Normal prior: $\Theta_3 \sim \text{Normal}(0, 1)$.

4. **Likelihood:**

$$Y_i \sim \text{Binomial}(1, p_i)$$

Here, the model predicts Y_i using a Binomial distribution with 1 trial (equivalent to a Bernoulli distribution). The logits calculated from the network are implicitly passed through the sigmoid function to represent the expected probability p_i . =====
In the Bayesian formulation, we place priors on all weights and biases and define a

likelihood for the output. For a classification task with a K -hidden-layer BNN with J neurons per hidden layer and a D_X -vector of predictors we can run the model as below. For the code example, we consider two hidden layers with a hyperbolic tangent (tanh) activation function, mapped to a logit output. Because the input matrix X incorporates the intercept as its first column, the bias term is implicitly included in the layer's weights:

$$\begin{aligned}
Y_i &\sim \text{Binomial}(1, p_i) \\
\text{logit}(p_i) &= H_{i,2} W_3 \\
H_{i,2} &= \tanh(H_{i,1} W_2) \\
H_{i,1} &= \tanh(X_i W_1) \\
W_1 &\sim \text{Normal}(0, 1) \\
W_2 &\sim \text{Normal}(0, 1) \\
W_3 &\sim \text{Normal}(0, 1)
\end{aligned}$$

where:

- Y_i is the observed binary outcome for the i -th observation.
- p_i is the predicted probability of the outcome being 1 for the i -th observation.
- X_i is the input row vector for the i -th observation, containing the intercept and the predictor variables. It has length $D_X = 2$.
- $H_{i,1}$ is the first hidden layer representation vector for the i -th observation. It has length $D_{H1} = 4$.
- $H_{i,2}$ is the second hidden layer representation vector for the i -th observation. It has length $D_{H2} = 3$.
- W_1 is the weight matrix of the first hidden layer, with a shape of $D_X \times D_{H1}$ (i.e., 2×4).
- W_2 is the weight matrix of the second hidden layer, with a shape of $D_{H1} \times D_{H2}$ (i.e., 4×3).
- W_3 is the final layer weight matrix used to compute the logits, with a shape of $D_{H2} \times 1$ (i.e., 3×1).
- All elements within the weight matrices W_1 , W_2 , and W_3 are assigned independent standard Normal priors. »»»> 20c35a832b7f5f82f2011e8aad8fe0f1e2851c1b

Notes

Note

- Using `m.dist.binomial(total_count=1)` alongside logit inputs is equivalent to specifying a standard binary cross-entropy loss with a sigmoid activation function in deep learning frameworks.
- The tutorial uses Stochastic Variational Inference (SVI) since exact MCMC ap-

proaches struggle with high-dimensional highly-correlated posterior spaces typical to neural networks. If you find MCMC chains not mixing during `m.fit()`, consider switching to an SVI backend or modifying priors to enforce further regularization.

Reference(s)

1. [PyData Berlin 2025: Introduction to Stochastic Variational Inference with NumPyro](#)
»»»< HEAD =====
2. [Variational Inference: Bayesian Neural Networks](#) »»»> 20c35a832b7f5f82f2011e8aad8fe0f1e2851c1b