

¹

Bayesian Inference

²

³ wordcount : 2955

⁴ **Abstract**

- ⁵ 1. Bayesian modeling is a powerful paradigm in modern statistics and machine learning, offering a principled framework for inference under uncertainty. However, practitioners face significant obstacles, including **interoperability** issues, a persistent **accessibility-flexibility trade-off**, the limitations of **domain-specific limitations**, and challenges in **scalability**.
- ¹⁰ 2. **Interoperability:** The landscape of Bayesian software is fragmented across programming languages and abstraction levels. Newcomers often gravitate towards high-level interfaces (e.g., *brms*) within familiar environments due to their accessibility for standard models. However, highgh levels of abstraction frameworks can be restrictive, lacking the flexibility needed for custom or complex models as research needs evolve.
- ¹⁵ 3. **Accessibility-Flexibility Trade-off:** To gain the necessary flexibility, researchers must often transition to lower-level probabilistic programming languages . This transition imposes a steeper learning curve and requires mastery of specific modeling languages or complex programming frameworks, hindering broader adoption and rapid iteration.
- ²⁰ 4. **Domain-Specific Limitations:** Similar accessibility and flexibility trade-offs exist in domain-specific Bayesian packages. While providing accessible, pre-packaged models for specific fields, customizing or extending these models often requires deep engagement with lower-level programming languages or switching tools entirely, limiting methodological innovation within those domains.
- ²⁵ 5. **Scalability:** Computational demands remain a significant bottleneck, limiting the application of Bayesian methods to the large datasets and complex, high-dimensional models prevalent in modern research.

28 6. To address these challenges, we introduce ***Bayesian Inference (BI)***, a new
29 Bayesian modeling software available in both Python and R. It aims to unify the
30 modeling experience by integrating an intuitive model-building syntax (enhancing
31 **accessibility**) with the **flexibility** of low-level abstraction coding available but also
32 pre-build function for high-level of abstraction and including hardware-accelerated
33 computation via JAX for improved **scalability**. Its availability in both major data
34 science languages directly tackles the **interoperability** barrier and the prebuild
35 function for specialized model in network analysis, survival models and phylogenetic
36 analysis allow to improved **domain-specific limitations**.

37 **Introduction**

38 Bayesian modeling has emerged as a vital tool in modern statistics and machine learning,
39 providing a framework for robust inference under uncertainty and the possibility to integrate
40 prior knowledge. Despite its potential, the practical application of Bayesian methods
41 is often hindered by significant hurdles within the current software ecosystem, preventing re-
42 searchers from fully leveraging its capabilities. Key challenges stem from the fragmented na-
43 ture of software across different programming languages (**interoperability**), gaps between
44 theoretical understanding and practical implementation (**accessibility**), complexities in
45 model specification that force trade-offs between ease-of-use and flexibility (**accessibility-**
46 **flexibility trade-off**), the constraints of overly specialized tools (**domain-specific limi-**
47 **tations**), and persistent computational scalability limitations for complex models or large
48 datasets (**scalability**).

49 The first major obstacle is the fragmented landscape of Bayesian software, scattered across
50 different programming languages and varying levels of abstraction, posing significant **in-**
51 **teroperability** challenges. Researchers frequently encounter a disparate collection of
52 tools—from *Stan*’s domain-specific language (DSL) to distinct low-level of abstraction
53 libraries (like *PyMC* (Salvatier, Wiecki, and Fonnesbeck 2016), *TensorFlow Probability*
54 (*TFP*) (Abadi et al. 2015), *NumPyro* (Phan, Pradhan, and Jankowiak 2019)) and high-
55 level of abstraction libraries (like *BRMS*). This fragmentation complicates workflows and
56 presents a confusing landscape, especially for researchers new to Bayesian analysis. For
57 instance, researchers new to Bayesian analysis may initially gravitate towards tools of
58 high-level of abstraction available within their most familiar programming environment
59 (e.g., *BRMS* (Bürkner 2017) in *R* (Wickham 2015)), potentially overlooking more suitable
60 options elsewhere due to the steep initial learning curve or perceived incompatibility (**ac-**
61 **cessibility**). This linguistic and platform diversity imposes considerable cognitive overhead,
62 potentially hindering the adoption of the most suitable tool for a given problem due to
63 familiarity biases or the friction of switching ecosystems, ultimately impacting the effec-
64 tive application of Bayesian methods. This initial hurdle of navigating disparate systems

65 naturally leads new practitioners to prioritize tools that appear easiest to learn, raising
66 concerns about the balance between accessibility and the flexibility needed for complex
67 research.

68 Compounding this fragmentation is the challenge of accessibility and the translation of
69 theoretical knowledge into practice **accessibility-flexibility trade-off**. Indeed, while
70 high-level interfaces like *BRMS* offer an intuitive formula-based syntax, significantly low-
71 ering the initial barrier to entry (e.g. generalized linear mixed models using *BRMS*), this
72 accessibility often comes at the cost of flexibility. As research questions become more
73 sophisticated, requiring custom likelihood functions (e.g., multiple likelihoods), intricate
74 prior structures (e.g., XXX), or non-standard model components (e.g., centered-random
75 factors), the limitations of these high-level wrappers become apparent. To gain the neces-
76 sary expressive power, the researcher must typically transition to lower-level probabilistic
77 programming languages (PPLs) such as *Stan* ([Stan Development Team](#)) (requiring mas-
78 terity of its specific DSL), *PyMC*, *NumPyro*, or *TFP*. This transition imposes a much steeper
79 learning curve, demanding a deeper understanding of probabilistic programming concepts
80 (like computational graphs or tensor manipulation) and often more verbose code. This
81 significant jump in complexity can deter users, divert focus from statistical modeling to
82 software engineering challenges, and ultimately slow down the pace of research, particularly
83 when trying to adapt models within specific scientific fields.

84 Similar accessibility and flexibility constraints manifest as **domain-specific limitations**
85 within specialized Bayesian packages. Fields like phylogenetics or network analysis benefit
86 from tools such as *BEAST* ([Bouckaert 2019](#)), *RevBayes* ([Höhna et al. 2016](#)), *STRAND*
87 ([Ross, McElreath, and Redhead 2024](#)), or *BISON* ([Hart et al. 2023](#)), which provide acces-
88 sible, pre-packaged models tailored to common domain problems. A phylogeneticist might
89 initially find *BEAST* convenient for standard molecular clock models. However, when
90 they wish to incorporate a novel evolutionary hypothesis requiring modification of the core
91 model structure or integrate data types not originally envisioned by the developers, they
92 often encounter rigid constraints. Extending these specialized tools frequently requires
93 deep engagement with their underlying, often complex, codebase (sometimes necessitat-
94 ing proficiency in languages like Java or C++) or abandoning the domain-specific tool
95 entirely in favor of a general-purpose PPL. This forces researchers to either compromise
96 on their methodological innovation or undertake a significant software development effort,
97 potentially switching programming ecosystems and losing the initial convenience, thereby
98 limiting the evolution of modeling practices within specialized domains. Even when model
99 specification is achievable, either in general or specialized tools, the computational feasi-
100 bility remains a major concern.

101 Finally, computational **scalability** continues to be a significant bottleneck, limiting the
102 application of Bayesian methods to the large datasets (e.g., millions of observations) and
103 complex, high-dimensional models (e.g., thousands of parameters) prevalent in modern re-

104 search across fields like genomics, neuroscience, and machine learning. While established
105 tools like *Stan* feature highly optimized inference algorithms (particularly its NUTS sam-
106 pler) and offer effective multi-core parallelization, they can still face challenges with long
107 C++ compilation times for complex models and may require substantial code restructur-
108 ing or external tooling to efficiently leverage hardware accelerators like GPUs or TPUs for
109 certain computations. Conversely, emerging frameworks built on *JAX* ([Bradbury et al.](#)
110 [2018](#)) (powering *NumPyro* and parts of *TFP*) promise substantial speedups via automatic
111 differentiation, JIT compilation, and native support for parallel hardware architectures.
112 However, integrating these powerful backends seamlessly into user-friendly, flexible model-
113 ing front-ends that don't require deep expertise in the JAX ecosystem itself is an ongoing
114 challenge. Domain-specific tools often inherit the scalability limitations of the frameworks
115 they are built upon, failing to provide a universally efficient solution across different model
116 types and data sizes.

117 Therefore, there is an evident and pressing demand for a Bayesian modeling framework
118 that synergistically addresses these interconnected limitations. To address these intercon-
119 nected challenges, we introduce ***Bayesian Inference (BI)***, a new Bayesian modeling
120 software designed to unify the modeling experience across the two dominant data science
121 languages, Python and R. *BI* tackles the **interoperability** barrier head-on by offering
122 native interfaces in both environments. It aims to resolve the **accessibility-flexibility**
123 **trade-off** by providing an intuitive model-building syntax familiar to users of statistical
124 modeling languages, while enabling advanced customization and leveraging multiple, inter-
125 changeable inference backends for flexibility. To combat **domain-specific limitations**,
126 *BI* includes pre-built functions and structures tailored for specialized models in areas like
127 network analysis, survival analysis, and phylogenetic analysis, while still allowing exten-
128 sion and modification within its general framework. Crucially, *BI* enhances **scalability**
129 by integrating with hardware-accelerated computation via *JAX* (using *NumPyro* or *TFP*
130 as backends), enabling efficient execution on CPUs, GPUs, and TPUs. By providing a
131 streamlined, efficient, and unified environment for the end-to-end Bayesian workflow—
132 from model specification and fitting to diagnostics and prediction—***BI*** lowers the barrier
133 to entry for sophisticated Bayesian modeling, aiming to empower a broader community
134 of researchers across disciplines to confidently apply advanced Bayesian methods to their
135 complex research problems.

136 Software Presentation

137 *BI* directly confronts the **interoperability** challenge by offering native, feature-equivalent
138 implementations in both Python and R. While minor syntactic differences exist to ad-
139 here to the idiomatic conventions of each language, the core model specification syntax,
140 the procedural workflow for analysis, and the underlying computational engines remain

141 fundamentally consistent. For instance, Python utilizes dot notation for method calls on
142 class objects (e.g., `bi.dist.normal(0,1)`), while R employs dollar sign notation for ac-
143 cessing elements or methods within its object system (e.g., `bi$dist$normal(0,1)`). This
144 dual-language availability significantly lowers the adoption barrier for researchers, allowing
145 them to work entirely within their preferred programming environment without sacrificing
146 access to a common, powerful Bayesian modeling framework.

147 *BI* is designed to navigate the critical *accessibility-flexibility trade-off* by providing multiple
148 layers of abstraction and utility, catering effectively to users with varying levels of Bayesian
149 modeling expertise and diverse complexity requirements through : simplified backend inter-
150 action via intuitive syntax, pre-built components for complex model features , addressing
151 domain-specific limitations within a general framework, integrated End-to-End Workflow
152 and extensive model library and documentation.

153 At its computational core, *BI* leverages the power and efficiency of established Probabilistic
154 Programming Languages (PPLs) like *NumPyro* and *TFP*, both of which are built upon the
155 *JAX* framework for high-performance numerical computation and automatic differentiation.
156 However, *BI* deliberately abstracts away much of the inherent complexity of these lower-
157 level tools (**Code block 1**). This significantly enhances **accessibility** for a broader range
158 of users.

159

160 **Code block 1:** Prior specification differences between *NumPyro*, *TFP*, and *BI*

```
162 # NumPyro prior specification
163 numpyro.sample("mu", dist.Normal(0, 1)).expand([10])

# TFP prior specification (within a JointDistributionCoroutine)
164 yield Root(tfd.Sample(tfd.Normal(loc=1.0, scale=1.0), sample_shape=10))

# BI prior specification
165 bi.dist.normal(0, 1, name = "mu", shape = (10,))
```

163 To enhance **flexibility** without unduly sacrificing the accessibility provided by the high-
164 level syntax, *BI* includes a library of pre-built, computationally optimized functions imple-
165 mented directly in *JAX* (e.g., **Code block 2**). These components encapsulate common but
166 potentially complex modeling structures, allowing users to incorporate them easily within
167 the model specification. Key examples include:

- 168 1. *Centered Random Effects* and *Non-Centered Random Effects* for hierarchical (multi-
169 level) model components (McElreath 2018). The non-centered parameterization, of-
170 ten crucial for efficient sampling in hierarchical models (particularly with sparse data),

171 is provided without requiring the user to manually implement the reparameterization
172 logic.

- 173 2. *Kernels for Gaussian Processes* for modeling spatial, temporal, phylogenetic, or other
174 forms of structured correlation or dependency.
- 175 3. *Block Model Effects* for implementing stochastic block models in network analysis.
- 176 4. *SRM effects* for modeling pairwise interactions in networks while accounting for
177 sender effects, receiver effects, dyadic effects, nodal predictors, dyadic predictors,
178 and observation biases ([Sosa et al., n.d.](#)).
- 179 5. *Network-Based Diffusion Approach (NBDA)* components for modeling the effect of
180 network edges on the rates of transmission of phenomena (e.g., behavioral, epidemi-
181 ological) while accounting for nodal or dyadic covariates.
- 182 6. *Network metrics* ranging from nodal, dyadic, and global network measures with a
183 total of 11 that can be used to build custom models of social network analysis ([Sosa,](#)
184 [Sueur, and Puga-Gonzalez 2020](#)).

185 These pre-built *JAX* functions provide tailored model components for common patterns in
186 specific fields, while keeping them fully integrated within the general, extensible modeling
187 framework. By providing these optimized building blocks within its general syntax, *BI*
188 allows researchers in these fields to rapidly implement standard domain models using fa-
189 miliar concepts. Crucially, however, users retain the full flexibility of the *BI* framework to
190 combine these domain-specific components with other model features (e.g., complex non-
191 linear effects via splines, hierarchical structures across groups of networks or phylogenies)
192 or to customize or extend them using *BI*'s underlying mechanisms if needed—a capabil-
193 ity often missing in more narrowly focused domain-specific packages. This design aims
194 to foster methodological innovation *within* specialized domains by lowering the barrier to
195 implementing more complex or novel models.

196

197 **Code block 2:** Example of random effect specification differences between NumPyro,
198 TFP, and BI

```
# NumPyro version of random centered effect
a = numpyro.sample("a", dist.Normal(5, 2))
    b = numpyro.sample("b", dist.Normal(-1, 0.5))
    sigma_cafe = numpyro.sample("sigma_cafe", dist.Exponential(1).expand([2]))
    sigma = numpyro.sample("sigma", dist.Exponential(1))
    Rho = numpyro.sample("Rho", dist.LKJ(2, 2))
    cov = jnp.outer(sigma_cafe, sigma_cafe) * Rho
    a_cafe_b_cafe = numpyro.sample(
```

```

    "a_cafe,b_cafe",
    dist.MultivariateNormal(jnp.stack([a, b]), cov).expand([20])
)
a_cafe, b_cafe = a_cafe_b_cafe[:, 0], a_cafe_b_cafe[:, 1]

# TFP version of random centered effect
alpha = yield Root(tfd.Sample(tfd.Normal(loc=5.0, scale=2.0), sample_shape=1))
beta = yield Root(tfd.Sample(tfd.Normal(loc=-1.0, scale=0.5), sample_shape=1))
sigma = yield Root(tfd.Sample(tfd.Exponential(rate=1.0), sample_shape=1))
sigma_alpha_beta = yield Root(tfd.Sample(tfd.Exponential(rate=1.0),
sample_shape=2))
Rho = yield Root(tfd.LKJ(dimension=2, concentration=2.0))
Mu = tf.concat([alpha, beta], axis=-1)
scale = tf.linalg.LinearOperatorDiag(sigma_alpha_beta).matmul(tf.squeeze(Rho))

# BI version of random centered effect
Sigma = dist.exponential(1, (ni,), name = 'Sigma_individual')
L = dist.lkjcholesky(1, (ni,), name = 'L_individual', shape = (ni,))
Z = dist.normal(0, 1, name = 'z_individual', shape = (ni,K))
alpha = random_centered2(Sigma, L, Z)

```

201 *BI* is designed to encapsulate the entire Bayesian modeling workflow within a cohesive
 202 object-oriented structure, promoting a streamlined and reproducible analysis pipeline. Typ-
 203 ically, a user interacts with a primary *BI* object, through which they can sequentially:

- 204 • **Handle Data:** Load, preprocess, and associate dataset(s) with the model object.
- 205 • **Define Model:** Specify the model structure, including the likelihood(s), priors for
 206 all parameters, and incorporate any pre-built components using an intuitive formula
 207 syntax.
- 208 • **Run Inference:** Execute the model fitting process using the No-U-Turn Sampler
 209 (NUTS), which triggers the backend PPL (e.g., *NumPyro*, *TFP*) to perform Markov
 210 Chain Monte Carlo (MCMC) sampling. Progress indicators and diagnostics are typ-
 211 ically provided.
- 212 • **Analyze Posterior:** Access, summarize, and diagnose the posterior distributions of
 213 parameters. This includes methods for calculating posterior means, medians, cred-
 214 ible intervals, convergence diagnostics (e.g., \hat{R} , Effective Sample Size - ESS), and
 215 retrieving raw posterior samples for custom analysis.

- 216 • **Visualize Results:** Generate standard diagnostic plots (e.g., trace plots, rank plots,
217 posterior distributions) and visualizations of model parameters, effects, and predic-
218 tions using integrated plotting functions that leverage the *arviz* library.

219 This unified structure minimizes the need for users to juggle multiple disparate software
220 tools or manually transfer data and results between different stages of the analysis, thereby
221 enhancing efficiency and reproducibility.

222 Finally, *BI* includes over 21 well-documented implementations of various standard and
223 advanced Bayesian models. Examples include Generalized Linear Models (GLMs), Gen-
224 eralized Linear Mixed Models (GLMMs), survival analysis models (e.g., Cox proportional
225 hazards), Principal Component Analysis (PCA), phylogenetic comparative methods, and
226 various network models. Each implementation is accompanied by detailed documentation
227 that encompasses: 1) general principles, 2) underlying assumptions, 3) code snippets in
228 Python and R, and 4) mathematical details, enabling users to gain a deeper understanding
229 of the modeling process and its nuances. Additionally, the framework's flexibility allows
230 models to be combined; for example, building a zero-inflated model with varying inter-
231 cepts and slopes, or constructing a joint model where principal components (derived from
232 PCA) serve as predictors in a subsequent regression, allowing uncertainty to be propagated
233 through all stages of the analysis.

234 **Example : SRM model**

235 To illustrate how these design features of *BI* coalesce to provide a streamlined, flexible, and
236 powerful solution, effectively addressing the limitations identified in the existing Bayesian
237 software landscape we will provide a basic example of how an SRM model is declared in
238 *BI*, compare it with the equivalent model in Numpyro (Appendix 1) and STAN (Appendix
239 2). We will also show how this model can be build from scratch with *BI* (**Code block**
240 **3**) or its custom functions (**Code block 4**) to highlight the accessibility-flexibility of our
241 package by demonstrating how advance user can build custom model (with less code than
242 STAN) as well as how new user can apply pre-build *BI* models. Finally we show how it
243 is also called in R (**Code block 5**) to cross language use with *BI*. Readers interested in
244 further details on data structure, data import, data manipulation, and model fitting for
245 SRM models can refer directly to the *BI* documentation [Modeling Network](#).

246
247 **Code block 3:** *SRM model from scratch with BI*
248

```
249 def model(N_id, idx, result_outcomes,  
250           focal_individual_predictors,  
251           target_individual_predictors):
```

```

# Intercept
intercept = bi.dist.normal(
    logit(0.1/jnp.sqrt(N_id)),
    2.5, shape=(1,), name = 'intercept'
)

# Sender receiver -----
N_var = focal_individual_predictors.shape[0]
N_id = focal_individual_predictors.shape[1]
focal_effects = dist.normal(0, 1, name = 'focal_effects')
target_effects = dist.normal( 0, 1, name = 'target_effects')
terms = jnp.stack([
    focal_effects @ focal_individual_predictors,
    target_effects @ target_individual_predictors
], axis = -1)
sr_raw = dist.normal(0, 1, shape=(2, N_id), name = 'sr_raw')
sr_sigma = dist.exponential( 1, shape= (2,), name = 'sr_sigma')
sr_L = dist.lkjcholesky(2, 2, name = "sr_L")
rf = deterministic('sr_rf',(((sr_L @ sr_raw).T * sr_sigma)))
ids = jnp.arange(0,sr_effects.shape[0])
edgl_idx = bi.net.vec_node_to_edgle(jnp.stack([ids, ids], axis = -1))
sender = sr_effects[edgl_idx[:,0],0] + sr_effects[edgl_idx[:,1],1]
receiver = sr_effects[edgl_idx[:,1],0] + sr_effects[edgl_idx[:,0],1]
sr = jnp.stack([sender, receiver], axis = 1)

# dyadic effects -----
bi.net.mat_to_edgl(dyadic_effect_mat)
dr_raw = dist.normal(0, 1, shape=(2,N_dyads), name = 'dr_raw')
dr_sigma = dist.exponential(1, name = 'dr_sigma' )
dr_L = dist.lkjcholesky(2, 2, name = 'dr_L')
dr_rf = deterministic('dr_rf', (
    ((dr_L @ dr_raw).T * jnp.repeat(dr_sigma, 2))
))

dyad_effects = dist.normal(0, 1,
    name= 'dyad_effects', shape = (dyadic_predictors.ndim - 1,
))
dr = dyad_effects * dyadic_predictors

# Likelihood

```

```
    bi.dist.poisson(jnp.exp(intercept + sr + dr), obs=result_outcomes)
```

250

Code block 4: SRM model with prebuild functions

```
251 def model(N_id, idx, result_outcomes,
252             focal_individual_predictors,
253             target_individual_predictors):
254
255     # Intercept
256     intercept = bi.dist.normal(
257         logit(0.1/jnp.sqrt(N_id)),
258         2.5, shape=(1,), name = 'intercept'
259     )
260
261     # SR
262     sr = bi.net.sender_receiver(
263         focal_individual_predictors,
264         target_individual_predictors
265     )
266
267     # Dyadic
268     dr = bi.net.dyadic_effect(shape = idx.shape[0])
269
270     # Likelihood
271     bi.dist.poisson(jnp.exp(intercept + sr + dr), obs=result_outcomes)
```

254

Code block 5: SRM model with prebuild functions

```
255 model <- function(N_id, idxShape, result_outcomes,
256                     focal_individual_predictors, target_individual_predictors){
257
258     x=0.1/jnp$sqrt(N_id)
259     tmp=jnp$log(x / (1 - x))
260
261     # Intercept
262     intercept = bi.dist.normal(tmp, 2.5, shape=c(1), name = 'block')
263
264     # SR
265     sr = m$net$sender_receiver(
```

```

        focal_individual_predictors,
        target_individual_predictors
    )

    # Dyadic
    dr = m$net$dyadic_effect(shape = c(idxShape))

    # Likelihood
    m$poisson(jnp$exp(intercept + sr + dr), obs=result_outcomes)
}

```

258 Finally, regarding code performance we can time the computation time for network of size
 259 200 in STAN and BI and observed that BI comput time is around XXX on cpu and XXX
 260 on gpu and STAN compute time around XXX.

261 Discussion

262 **BI** framework is built on top of the popular Python programming language, with a fo-
 263 cus on providing a user-friendly interface for model development and interpretation. Our
 264 framework is designed to be modular and extensible, allowing users to easily incorporate
 265 their own custom models and data types into the framework. One of the key features of
 266 this software is its comprehensive library of 21 predefined Bayesian models, covering a wide
 267 range of common applications and use cases. These models are accompanied by detailed
 268 explanations, making it easier for users to understand the underlying assumptions and
 269 apply the models to their specific research questions. In addition to these built-in models,
 270 the software includes several custom functions tailored for advanced statistical and network
 271 modeling. This curated library serves not only as a collection of ready-to-use tools but also
 272 as a valuable pedagogical resource, demonstrating best practices for constructing, fitting,
 273 and interpreting models within the *BI* framework, and providing robust templates for users
 274 aiming to develop novel model variants. Whether users are interested in hierarchical mod-
 275 els, time-series analysis, or cutting-edge network modeling approaches, our library caters
 276 to a variety of analytical needs. This accessibility fosters an environment where users can
 277 confidently explore and implement Bayesian methods, ultimately enhancing their research
 278 capabilities.

279 By providing a streamlined and efficient environment for the end-to-end Bayesian workflow—
 280 from model specification and fitting to diagnostics and prediction, *BI* lowers the barrier
 281 to entry for sophisticated Bayesian modeling. We aim to empower a broader community
 282 of researchers across disciplines to confidently apply advanced Bayesian methods to their
 283 complex research problems.

284 **References**

- 285 Abadi, Martín, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro,
286 Greg S. Corrado, et al. 2015. “TensorFlow: Large-Scale Machine Learning on Hetero-
287 geneous Systems.” <https://www.tensorflow.org/>.
- 288 Bouckaert, Timothy G. AND Barido-Sottani, Remco AND Vaughan. 2019. “BEAST 2.5:
289 An Advanced Software Platform for Bayesian Evolutionary Analysis.” *PLOS Compu-
290 tational Biology* 15 (4): 1–28. <https://doi.org/10.1371/journal.pcbi.1006650>.
- 291 Bradbury, James, Roy Frostig, Peter Hawkins, Matthew James Johnson, Chris Leary,
292 Dougal Maclaurin, George Necula, et al. 2018. “JAX: Composable Transformations of
293 Python+NumPy Programs.” <http://github.com/jax-ml/jax>.
- 294 Bürkner, Paul-Christian. 2017. “brms: An R Package for Bayesian Multilevel Models
295 Using Stan.” *Journal of Statistical Software* 80 (1): 1–28. <https://doi.org/10.18637/jss.v080.i01>.
- 296 Hart, Jordan, Michael Nash Weiss, Daniel Franks, and Lauren Brent. 2023. “BISoN:
297 A Bayesian Framework for Inference of Social Networks.” *Methods in Ecology and
298 Evolution* 14 (9): 2411–20. <https://doi.org/10.1111/2041-210X.14171>.
- 300 Höhna, Sebastian, Michael J. Landis, Tracy A. Heath, Bastien Boussau, Nicolas Lartillot,
301 Brian R. Moore, John P. Huelsenbeck, and Fredrik Ronquist. 2016. “RevBayes:
302 Bayesian Phylogenetic Inference Using Graphical Models and an Interactive Model-
303 Specification Language.” *Systematic Biology* 65 (4): 726–36. <https://doi.org/10.1093/sysbio/syw021>.
- 305 McElreath, Richard. 2018. *Statistical Rethinking: A Bayesian Course with Examples in r
306 and Stan*. Chapman; Hall/CRC.
- 307 Phan, Du, Neeraj Pradhan, and Martin Jankowiak. 2019. “Composable Effects for
308 Flexible and Accelerated Probabilistic Programming in NumPyro.” *arXiv Preprint
309 arXiv:1912.11554*.
- 310 Ross, Cody T, Richard McElreath, and Daniel Redhead. 2024. “Modelling Animal Network
311 Data in r Using STRAND.” *Journal of Animal Ecology* 93 (3): 254–66.
- 312 Salvatier, John, Thomas V Wiecki, and Christopher Fonnesbeck. 2016. “Probabilistic
313 Programming in Python Using PyMC3.” *PeerJ Computer Science* 2: e55.
- 314 Sosa, Sebastian, Mary B. McElreath, Daniel Redhead, and Cody T. Ross. n.d. “Robust
315 Bayesian Analysis of Animal Networks Subject to Biases in Sampling Intensity and
316 Censoring.” *Methods in Ecology and Evolution* n/a (n/a). <https://doi.org/https://doi.org/10.1111/2041-210X.70017>.
- 318 Sosa, Sebastian, Cédric Sueur, and Ivan Puga-Gonzalez. 2020. “Network Measures in
319 Animal Social Network Analysis: Their Strengths, Limits, Interpretations and Uses.”
- 320 Stan Development Team. “Stan Modeling Language Users Guide and Reference Manual,
321 Version 2.32.” <https://mc-stan.org>.
- 322 Wickham, Hadley. 2015. *R Packages*. 1st ed. O’Reilly Media, Inc.