

# Bayesian Inference library

3 Sebastian Sosa<sup>1\*</sup>, Mary Brooke McElreath<sup>1</sup>, Cody T. Ross<sup>1</sup>

7 wordcount : 2959

## 8 Abstract

- 9 1. Bayesian modeling is a powerful paradigm in modern statistics and machine learning, offering a principled  
10 framework for inference under uncertainty. However, practitioners face significant obstacles, including  
11 **interoperability** issues, a persistent **accessibility-flexibility trade-off**, the limitations of **domain-**  
12 **specific limitations**, and challenges in **scalability**.
- 13 2. **Interoperability:** The landscape of Bayesian software is fragmented across programming languages and  
14 abstraction levels. Newcomers often gravitate towards high-level interfaces (e.g., *brms*) within familiar  
15 environments due to their accessibility for standard models. However, highgh levels of abstraction frame-  
16 works can be restrictive, lacking the flexibility needed for custom or complex models as research needs  
17 evolve.
- 18 3. **Accessibility-Flexibility Trade-off:** To gain the necessary flexibility, researchers must often transition  
19 to lower-level probabilistic programming languages . This transition imposes a steeper learning curve and  
20 requires mastery of specific modeling languages or complex programming frameworks, hindering broader  
21 adoption and rapid iteration.
- 22 4. **Domain-Specific Limitations:** Similar accessibility and flexibility trade-offs exist in domain-specific  
23 Bayesian packages. While providing accessible, pre-packaged models for specific fields, customizing or  
24 extending these models often requires deep engagement with lower-level programming languages or switch-  
25 ing tools entirely, limiting methodological innovation within those domains.
- 26 5. **Scalability:** Computational demands remain a significant bottleneck, limiting the application of Bayesian  
27 methods to the large datasets and complex, high-dimensional models prevalent in modern research.
- 28 6. To address these challenges, we introduce ***Bayesian Inference (BI)***, a new Bayesian modeling software  
29 available in both Python and R. It aims to unify the modeling experience by integrating an intuitive model-  
30 building syntax (enhancing **accessibility**) with the **flexibility** of low-level abstraction coding available  
31 but also pre-build function for high-level of abstraction and including hardware-accelerated computation  
32 via JAX for improved **scalability**. Its availability in both major data science languages directly tackles  
33 the **interoperability** barrier and the prebuild function for specialized model in network analysis, survival  
34 models and phylogenetic analysis allow to improved **domain-specific limitations**.

## 35 Introduction

36 Bayesian modeling has emerged as a vital tool in modern statistics and machine learning, providing a framework  
37 for robust inference under uncertainty and the possibility to integrate prior knowledge. Despite its potential,  
38 the practical application of Bayesian methods is often hindered by significant hurdles within the current  
39 software ecosystem, preventing researchers from fully leveraging its capabilities. Key challenges stem from  
40 the fragmented nature of software across different programming languages (**interoperability**), gaps between  
41 theoretical understanding and practical implementation (**accessibility**), complexities in model specification  
42 that force trade-offs between ease-of-use and flexibility (**accessibility-flexibility trade-off**), the constraints

43 of overly specialized tools (**domain-specific limitations**), and persistent computational scalability limitations  
44 for complex models or large datasets (**scalability**).

45 The first major obstacle is the fragmented landscape of Bayesian software, scattered across different program-  
46 ming languages and varying levels of abstraction, posing significant **interoperability** challenges. Researchers  
47 frequently encounter a disparate collection of tools—from *Stan*'s domain-specific language (DSL) to distinct  
48 low-level abstraction libraries (like *PyMC* (Salvatier, Wiecki, and Fonnesbeck 2016), *TensorFlow Probability*  
49 (*TFP*) (Abadi et al. 2015), *NumPyro* (Phan, Pradhan, and Jankowiak 2019)) and high-level abstraction  
50 libraries (like *BRMS*). This fragmentation complicates workflows and presents a confusing landscape, especially  
51 for researchers new to Bayesian analysis. For instance, researchers new to Bayesian analysis may initially  
52 gravitate towards tools of high-level abstraction available within their most familiar programming environment  
53 (e.g., *BRMS* (Bürkner 2017) in *R* (Wickham 2015)), potentially overlooking more suitable options elsewhere  
54 due to the steep initial learning curve or perceived incompatibility (accessibility). This linguistic and platform  
55 diversity imposes considerable cognitive overhead, potentially hindering the adoption of the most suitable tool  
56 for a given problem due to familiarity biases or the friction of switching ecosystems, ultimately impacting the  
57 effective application of Bayesian methods. This initial hurdle of navigating disparate systems naturally leads  
58 new practitioners to prioritize tools that appear easiest to learn, raising concerns about the balance between  
59 accessibility and the flexibility needed for complex research.

60 Compounding this fragmentation is the challenge of accessibility and the translation of theoretical knowledge  
61 into practice **accessibility-flexibility trade-off**. Indeed, while high-level interfaces like *BRMS* offer an intuitive  
62 formula-based syntax, significantly lowering the initial barrier to entry (e.g. generalized linear mixed models  
63 using *BRMS*), this accessibility often comes at the cost of flexibility. As research questions become more  
64 sophisticated, requiring custom likelihood functions (e.g., multiple likelihoods), intricate prior structures (e.g.,  
65 XXX), or non-standard model components (e.g., centered-random factors), the limitations of these high-level  
66 wrappers become apparent. To gain the necessary expressive power, the researcher must typically transition  
67 to lower-level probabilistic programming languages (PPLs) such as *Stan* (Stan Development Team) (requiring  
68 mastery of its specific DSL), *PyMC*, *NumPyro*, or *TFP*. This transition imposes a much steeper learning curve,  
69 demanding a deeper understanding of probabilistic programming concepts (like computational graphs or tensor  
70 manipulation) and often more verbose code. This significant jump in complexity can deter users, divert focus  
71 from statistical modeling to software engineering challenges, and ultimately slow down the pace of research,  
72 particularly when trying to adapt models within specific scientific fields.

73 Similar accessibility and flexibility constraints manifest as **domain-specific limitations** within specialized  
74 Bayesian packages. Fields like phylogenetics or network analysis benefit from tools such as *BEAST* (Bouckaert  
75 2019), *RevBayes* (Höhna et al. 2016), *STRAND* (Ross, McElreath, and Redhead 2024), or *BISON* (Hart et al.  
76 2023), which provide accessible, pre-packaged models tailored to common domain problems. A phylogeneticist  
77 might initially find *BEAST* convenient for standard molecular clock models. However, when they wish to  
78 incorporate a novel evolutionary hypothesis requiring modification of the core model structure or integrate  
79 data types not originally envisioned by the developers, they often encounter rigid constraints. Extending these  
80 specialized tools frequently requires deep engagement with their underlying, often complex, codebase (sometimes  
81 necessitating proficiency in languages like Java or C++) or abandoning the domain-specific tool entirely in favor  
82 of a general-purpose PPL. This forces researchers to either compromise on their methodological innovation or  
83 undertake a significant software development effort, potentially switching programming ecosystems and losing  
84 the initial convenience, thereby limiting the evolution of modeling practices within specialized domains. Even  
85 when model specification is achievable, either in general or specialized tools, the computational feasibility  
86 remains a major concern.

87 Finally, computational **scalability** continues to be a significant bottleneck, limiting the application of Bayesian  
88 methods to the large datasets (e.g., millions of observations) and complex, high-dimensional models (e.g.,  
89 thousands of parameters) prevalent in modern research across fields like genomics, neuroscience, and machine  
90 learning. While established tools like *Stan* feature highly optimized inference algorithms (particularly its NUTS  
91 sampler) and offer effective multi-core parallelization, they can still face challenges with long C++ compilation  
92 times for complex models and may require substantial code restructuring or external tooling to efficiently  
93 leverage hardware accelerators like GPUs or TPUs for certain computations. Conversely, emerging frameworks  
94 built on *JAX* (Bradbury et al. 2018) (powering *NumPyro* and parts of *TFP*) promise substantial speedups  
95 via automatic differentiation, JIT compilation, and native support for parallel hardware architectures. However,  
96 integrating these powerful backends seamlessly into user-friendly, flexible modeling front-ends that don't require  
97 deep expertise in the *JAX* ecosystem itself is an ongoing challenge. Domain-specific tools often inherit the  
98 scalability limitations of the frameworks they are built upon, failing to provide a universally efficient solution  
99 across different model types and data sizes.

100 Therefore, there is an evident and pressing demand for a Bayesian modeling framework that synergistically  
 101 addresses these interconnected limitations. To address these interconnected challenges, we introduce ***Bayesian***  
 102 ***Inference (BI)***, a new Bayesian modeling software designed to unify the modeling experience across the  
 103 two dominant data science languages, Python and R. ***BI*** tackles the **interoperability** barrier head-on by  
 104 offering native interfaces in both environments. It aims to resolve the **accessibility-flexibility trade-off** by  
 105 providing an intuitive model-building syntax familiar to users of statistical modeling languages, while enabling  
 106 advanced customization and leveraging multiple, interchangeable inference backends for flexibility. To combat  
 107 **domain-specific limitations**, ***BI*** includes pre-built functions and structures tailored for specialized models  
 108 in areas like network analysis, survival analysis, and phylogenetic analysis, while still allowing extension and  
 109 modification within its general framework. Crucially, ***BI*** enhances **scalability** by integrating with hardware-  
 110 accelerated computation via *JAX* (using *NumPyro* or *TFP* as backends), enabling efficient execution on CPUs,  
 111 GPUs, and TPUs. By providing a streamlined, efficient, and unified environment for the end-to-end Bayesian  
 112 workflow—from model specification and fitting to diagnostics and prediction—***BI*** lowers the barrier to entry  
 113 for sophisticated Bayesian modeling, aiming to empower a broader community of researchers across disciplines  
 114 to confidently apply advanced Bayesian methods to their complex research problems.

## 115 Software Presentation

116 ***BI*** directly confronts the **interoperability** challenge by offering native, feature-equivalent implementations  
 117 in both Python and R. While minor syntactic differences exist to adhere to the idiomatic conventions of  
 118 each language, the core model specification syntax, the procedural workflow for analysis, and the underlying  
 119 computational engines remain fundamentally consistent. For instance, Python utilizes dot notation for method  
 120 calls on class objects (e.g., `bi.dist.normal(0,1)`), while R employs dollar sign notation for accessing elements  
 121 or methods within its object system (e.g., `bi$dist$normal(0,1)`). This dual-language availability significantly  
 122 lowers the adoption barrier for researchers, allowing them to work entirely within their preferred programming  
 123 environment without sacrificing access to a common, powerful Bayesian modeling framework.

124 ***BI*** is designed to navigate the critical **accessibility-flexibility trade-off** by providing multiple layers of abstraction  
 125 and utility, catering effectively to users with varying levels of Bayesian modeling expertise and diverse complexity  
 126 requirements through : simplified backend interaction via intuitive syntax, pre-built components for complex  
 127 model features , addressing domain-specific limitations within a general framework, integrated End-to-End  
 128 Workflow and extensive model library and documentation.

129 At its computational core, ***BI*** leverages the power and efficiency of established Probabilistic Programming Lan-  
 130 guages (PPLs) like *NumPyro* and *TFP*, both of which are built upon the *JAX* framework for high-performance  
 131 numerical computation and automatic differentiation. However, ***BI*** deliberately abstracts away much of the  
 132 inherent complexity of these lower-level tools (**Code block 1**). This significantly enhances **accessibility** for a  
 133 broader range of users.

---

134  
 135 ***Code block 1:*** Prior specification differences between *NumPyro*, *TFP*, and ***BI***

---

```
# NumPyro prior specification
numppyro.sample("mu", dist.Normal(0, 1)).expand([10])

# TFP prior specification (within a JointDistributionCoroutine)
yield Root(tfd.Sample(tfd.Normal(loc=1.0, scale=1.0), sample_shape=10))

# BI prior specification
bi.dist.normal(0, 1, name = "mu", shape = (10,))
```

138 To enhance **flexibility** without unduly sacrificing the accessibility provided by the high-level syntax, ***BI*** includes  
 139 a library of pre-built, computationally optimized functions implemented directly in *JAX* (e.g., **Code block**  
 140 **2**). These components encapsulate common but potentially complex modeling structures, allowing users to  
 141 incorporate them easily within the model specification. Key examples include:

- 142 1. *Centered Random Effects* and *Non-Centered Random Effects* for hierarchical (multi-level) model com-  
 143 ponents (McElreath 2018). The non-centered parameterization, often crucial for efficient sampling in  
 144 hierarchical models (particularly with sparse data), is provided without requiring the user to manually  
 145 implement the reparameterization logic.
- 146 2. *Kernels for Gaussian Processes* for modeling spatial, temporal, phylogenetic, or other forms of structured  
 147 correlation or dependency.
- 148 3. *Block Model Effects* for implementing stochastic block models in network analysis.

- 149 4. *SRM effects* for modeling pairwise interactions in networks while accounting for sender effects, receiver  
 150 effects, dyadic effects, nodal predictors, dyadic predictors, and observation biases (Sosa et al., n.d.).  
 151 5. *Network-Based Diffusion Approach (NBDA)* components for modeling the effect of network edges on the  
 152 rates of transmission of phenomena (e.g., behavioral, epidemiological) while accounting for nodal or dyadic  
 153 covariates.  
 154 6. *Network metrics* ranging from nodal, dyadic, and global network measures with a total of 11 that can be  
 155 used to build custom models of social network analysis (Sosa, Sueur, and Puga-Gonzalez 2020).

156 These pre-built *JAX* functions provide tailored model components for common patterns in specific fields, while  
 157 keeping them fully integrated within the general, extensible modeling framework. By providing these optimized  
 158 building blocks within its general syntax, *BI* allows researchers in these fields to rapidly implement standard  
 159 domain models using familiar concepts. Crucially, however, users retain the full flexibility of the *BI* framework  
 160 to combine these domain-specific components with other model features (e.g., complex non-linear effects via  
 161 splines, hierarchical structures across groups of networks or phylogenies) or to customize or extend them using  
 162 *BI*'s underlying mechanisms if needed—a capability often missing in more narrowly focused domain-specific  
 163 packages. This design aims to foster methodological innovation *within* specialized domains by lowering the  
 164 barrier to implementing more complex or novel models [link to latex block].

165 **Code block 2:** Random effect specification differences between NumPyro, TFP, and BI

```
166 # Numpyro version of random centered effect
167 a = numpyro.sample("a", dist.Normal(5, 2))
  b = numpyro.sample("b", dist.Normal(-1, 0.5))
  sigma_cafe = numpyro.sample("sigma_cafe", dist.Exponential(1).expand([2]))
  sigma = numpyro.sample("sigma", dist.Exponential(1))
  Rho = numpyro.sample("Rho", dist.LKJ(2, 2))
  cov = jnp.outer(sigma_cafe, sigma_cafe) * Rho
  a_cafe_b_cafe = numpyro.sample(
    "a_cafe,b_cafe",
    dist.MultivariateNormal(jnp.stack([a, b]), cov).expand([20])
  )
a_cafe, b_cafe = a_cafe_b_cafe[:, 0], a_cafe_b_cafe[:, 1]

# TFP version of random centered effect
alpha = yield Root(tfd.Sample(tfd.Normal(loc=5.0, scale=2.0), sample_shape=1))
beta = yield Root(tfd.Sample(tfd.Normal(loc=-1.0, scale=0.5), sample_shape=1))
sigma = yield Root(tfd.Sample(tfd.Exponential(rate=1.0), sample_shape=1))
sigma_alpha_beta = yield Root(tfd.Sample(tfd.Exponential(rate=1.0),
sample_shape=2))
Rho = yield Root(tfd.LKJ(dimension=2, concentration=2.0))
Mu = tf.concat([alpha, beta], axis=-1)
scale = tf.linalg.LinearOperatorDiag(sigma_alpha_beta).matmul(tf.squeeze(Rho))

# BI version of random centered effect
Sigma = dist.exponential(1, (ni,), name = 'Sigma_individual')
L = dist.lkjcholesky(1, (ni,), name = 'L_individual', shape = (ni,))
Z = dist.normal(0, 1, name = 'z_individual', shape = (ni,K))
alpha = random_centered2(Sigma, L, Z)
```

169 *BI* is designed to encapsulate the entire Bayesian modeling workflow within a cohesive object-oriented structure,  
 170 promoting a streamlined and reproducible analysis pipeline. Typically, a user interacts with a primary *BI* object,  
 171 through which they can sequentially:

- 172 • **Handle Data:** Load, preprocess, and associate dataset(s) with the model object.
- 173 • **Define Model:** Specify the model structure, including the likelihood(s), priors for all parameters, and  
 174 incorporate any pre-built components using an intuitive formula syntax.
- 175 • **Run Inference:** Execute the model fitting process using the No-U-Turn Sampler (NUTS), which triggers  
 176 the backend PPL (e.g., *NumPyro*, *TFP*) to perform Markov Chain Monte Carlo (MCMC) sampling.  
 177 Progress indicators and diagnostics are typically provided.

- **Analyze Posterior:** Access, summarize, and diagnose the posterior distributions of parameters. This includes methods for calculating posterior means, medians, credible intervals, convergence diagnostics (e.g.,  $\hat{R}$ , Effective Sample Size - ESS), and retrieving raw posterior samples for custom analysis.
- **Visualize Results:** Generate standard diagnostic plots (e.g., trace plots, rank plots, posterior distributions) and visualizations of model parameters, effects, and predictions using integrated plotting functions that leverage the *arviz* library.

This unified structure minimizes the need for users to juggle multiple disparate software tools or manually transfer data and results between different stages of the analysis, thereby enhancing efficiency and reproducibility.

Finally, *BI* includes over 21 well-documented implementations of various standard and advanced Bayesian models. Examples include Generalized Linear Models (GLMs), Generalized Linear Mixed Models (GLMMs), survival analysis models (e.g., Cox proportional hazards), Principal Component Analysis (PCA), phylogenetic comparative methods, and various network models. Each implementation is accompanied by detailed documentation that encompasses: 1) general principles, 2) underlying assumptions, 3) code snippets in Python and R, and 4) mathematical details, enabling users to gain a deeper understanding of the modeling process and its nuances. Additionally, the framework's flexibility allows models to be combined; for example, building a zero-inflated model with varying intercepts and slopes, or constructing a joint model where principal components (derived from PCA) serve as predictors in a subsequent regression, allowing uncertainty to be propagated through all stages of the analysis.

## 196 Example : SRM model

To illustrate how these design features of *BI* coalesce to provide a streamlined, flexible, and powerful solution, effectively addressing the limitations identified in the existing Bayesian software landscape we will provide a basic example of how an SRM model is declared in *BI*, compare it with the equivalent model in Numpyro (Appendix 1) and STAN (Appendix 2). We will also show how this model can be build from scratch with *BI* (**Code block 3**) or its custom functions (**Code block 4**) to highlight the accessibility-flexibility of our package by demonstrating how advance user can build custom model (with less code than STAN) as well as how new user can apply pre-build *BI* models. Finally we show how it is also called in R (**Code block 5**) to cross language use with *BI*. Readers interested in further details on data structure, data import, data manipulation, and model fitting for SRM models can refer directly to the *BI* documentation [Modeling Network](#).

---

### 206 *Code block 3: SRM model from scratch with BI*

```
207 def model(N_id, idx, result_outcomes,
208     focal_individual_predictors,
209     target_individual_predictors):
210
211     # Intercept
212     intercept = bi.dist.normal(
213         logit(0.1/jnp.sqrt(N_id)),
214         2.5, shape=(1,), name = 'intercept'
215     )
216
217     # Sender receiver -----
218     N_var = focal_individual_predictors.shape[0]
219     N_id = focal_individual_predictors.shape[1]
220     focal_effects = dist.normal(0, 1, name = 'focal_effects')
221     target_effects = dist.normal( 0, 1, name = 'target_effects')
222     terms = jnp.stack([
223         focal_effects @ focal_individual_predictors,
224         target_effects @ target_individual_predictors
225     ], axis = -1)
226
227     sr_raw = dist.normal(0, 1, shape=(2, N_id), name = 'sr_raw')
228     sr_sigma = dist.exponential( 1, shape= (2,), name = 'sr_sigma')
229     sr_L = dist.lkjcholesky(2, 2, name = "sr_L")
230     rf = deterministic('sr_rf',(((sr_L @ sr_raw).T * sr_sigma)))
231     ids = jnp.arange(0,sr_effects.shape[0])
232     edgl_idx = bi.net.vec_node_to_edg(e(jnp.stack([ids, ids], axis = -1)))
233     sender = sr_effects[edgl_idx[:,0],0] + sr_effects[edgl_idx[:,1],1]
```

```

receiver = sr_effects[edgl_idx[:,1],0] + sr_effects[edgl_idx[:,0],1]
sr = jnp.stack([sender, receiver], axis = 1)

# dyadic effects -----
bi.net.mat_to_edgl(dyadic_effect_mat)
dr_raw = dist.normal(0, 1, shape=(2,N_dyads), name = 'dr_raw')
dr_sigma = dist.exponential(1, name = 'dr_sigma' )
dr_L = dist.lkjcholesky(2, 2, name = 'dr_L')
dr_rf = deterministic('dr_rf', (
    ((dr_L @ dr_raw).T * jnp.repeat(dr_sigma, 2))
))

dyad_effects = dist.normal(0, 1,
    name= 'dyad_effects', shape = (dyadic_predictors.ndim - 1,
))
dr = dyad_effects * dyadic_predictors

# Likelihood
bi.dist.poisson(jnp.exp(intercept + sr + dr), obs=result_outcomes)

```

210

211

**Code block 4:** SRM model with prebuild functions

```

def model(N_id, idx, result_outcomes,
          focal_individual_predictors,
          target_individual_predictors):

    # Intercept
    intercept = bi.dist.normal(
        logit(0.1/jnp.sqrt(N_id)),
        2.5, shape=(1,), name = 'intercept'
    )

    # SR
    sr = bi.net.sender_receiver(
        focal_individual_predictors,
        target_individual_predictors
    )

    # Dyadic
    dr = bi.net.dyadic_effect(shape = idx.shape[0])

    # Likelihood
    bi.dist.poisson(jnp.exp(intercept + sr + dr), obs=result_outcomes)

```

214

215

**Code block 5:** SRM model with prebuild functions

```

model <- function(N_id, idxShape, result_outcomes,
                   focal_individual_predictors, target_individual_predictors){

  x=0.1/jnp$sqrt(N_id)
  tmp=jnp$log(x / (1 - x))

  # Intercept
  intercept = bi.dist.normal(tmp, 2.5, shape=c(1), name = 'block')

  # SR
  sr = m$net$sender_receiver(
      focal_individual_predictors,
      target_individual_predictors

```

```

    )

# Dyadic
dr = m$net$dyadic_effect(shape = c(idxShape))

# Likelihood
m$poisson(jnp$exp(intercept + sr + dr), obs=result_outcomes)
}

```

218 Finally, regarding code performance we can time the computation time for network of size 200 in STAN and BI  
 219 and observed that BI comput time is around XXX on cpu and XXX on gpu and STAN compute time around  
 220 XXX.

## 221 Discussion

222 BI framework is built on top of the popular Python programming language, with a focus on providing a user-  
 223 friendly interface for model development and interpretation. Our framework is designed to be modular and  
 224 extensible, allowing users to easily incorporate their own custom models and data types into the framework.  
 225 One of the key features of this software is its comprehensive library of 21 predefined Bayesian models, covering  
 226 a wide range of common applications and use cases. These models are accompanied by detailed explanations,  
 227 making it easier for users to understand the underlying assumptions and apply the models to their specific  
 228 research questions. In addition to these built-in models, the software includes several custom functions tailored  
 229 for advanced statistical and network modeling. This curated library serves not only as a collection of ready-  
 230 to-use tools but also as a valuable pedagogical resource, demonstrating best practices for constructing, fitting,  
 231 and interpreting models within the BI framework, and providing robust templates for users aiming to develop  
 232 novel model variants. Whether users are interested in hierarchical models, time-series analysis, or cutting-edge  
 233 network modeling approaches, our library caters to a variety of analytical needs. This accessibility fosters an  
 234 environment where users can confidently explore and implement Bayesian methods, ultimately enhancing their  
 235 research capabilities.

236 By providing a streamlined and efficient environment for the end-to-end Bayesian workflow—from model  
 237 specification and fitting to diagnostics and prediction, BI lowers the barrier to entry for sophisticated Bayesian  
 238 modeling. We aim to empower a broader community of researchers across disciplines to confidently apply  
 239 advanced Bayesian methods to their complex research problems.

## 240 References

- 241 Abadi, Martín, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, Greg S. Corrado, et al.  
 242 2015. “TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems.” <https://www.tensorflow.org/>.  
 243 Bouckaert, Timothy G. AND Barido-Sottani, Remco AND Vaughan. 2019. “BEAST 2.5: An Advanced Software  
 244 Platform for Bayesian Evolutionary Analysis.” *PLOS Computational Biology* 15 (4): 1–28. <https://doi.org/10.1371/journal.pcbi.1006650>.  
 245 Bradbury, James, Roy Frostig, Peter Hawkins, Matthew James Johnson, Chris Leary, Dougal Maclaurin, George  
 246 Necula, et al. 2018. “JAX: Composable Transformations of Python+NumPy Programs.” <http://github.com/jax-ml/jax>.  
 247 Bürkner, Paul-Christian. 2017. “brms: An R Package for Bayesian Multilevel Models Using Stan.” *Journal of  
 248 Statistical Software* 80 (1): 1–28. <https://doi.org/10.18637/jss.v080.i01>.  
 249 Hart, Jordan, Michael Nash Weiss, Daniel Franks, and Lauren Brent. 2023. “BISON: A Bayesian Framework  
 250 for Inference of Social Networks.” *Methods in Ecology and Evolution* 14 (9): 2411–20. <https://doi.org/https://doi.org/10.1111/2041-210X.14171>.  
 251 Höhna, Sebastian, Michael J. Landis, Tracy A. Heath, Bastien Boussau, Nicolas Lartillot, Brian R. Moore, John  
 252 P. Huelsenbeck, and Fredrik Ronquist. 2016. “RevBayes: Bayesian Phylogenetic Inference Using Graphical  
 253 Models and an Interactive Model-Specification Language.” *Systematic Biology* 65 (4): 726–36. <https://doi.org/10.1093/sysbio/syw021>.  
 254 McElreath, Richard. 2018. *Statistical Rethinking: A Bayesian Course with Examples in r and Stan*. Chapman;  
 255 Hall/CRC.  
 256 Phan, Du, Neeraj Pradhan, and Martin Jankowiak. 2019. “Composable Effects for Flexible and Accelerated  
 257 Probabilistic Programming in NumPyro.” *arXiv Preprint arXiv:1912.11554*.  
 258 Ross, Cody T, Richard McElreath, and Daniel Redhead. 2024. “Modelling Animal Network Data in r Using  
 259 STRAND.” *Journal of Animal Ecology* 93 (3): 254–66.  
 260 Salvatier, John, Thomas V Wiecki, and Christopher Fonnesbeck. 2016. “Probabilistic Programming in Python  
 261 Using PyMC3.” *PeerJ Computer Science* 2: e55.

- 266 Sosa, Sebastian, Mary B. McElreath, Daniel Redhead, and Cody T. Ross. n.d. "Robust Bayesian Analysis of  
267 Animal Networks Subject to Biases in Sampling Intensity and Censoring." *Methods in Ecology and Evolution*  
268 n/a (n/a). <https://doi.org/10.1111/2041-210X.70017>.
- 269 Sosa, Sebastian, Cédric Sueur, and Ivan Puga-Gonzalez. 2020. "Network Measures in Animal Social Network  
270 Analysis: Their Strengths, Limits, Interpretations and Uses."
- 271 Stan Development Team. "Stan Modeling Language Users Guide and Reference Manual, Version 2.32." <https://mc-stan.org>.
- 272 Wickham, Hadley. 2015. *R Packages*. 1st ed. O'Reilly Media, Inc.