

¹

Bayesian Inference library

²

Sebastian Sosa^{1,*}

³

Mary B. McElreath¹

⁴

Cody T. Ross¹

⁵ ¹ Department of Human Behaviour, Ecology and Culture, Max Planck Institute for Evolutionary Anthropology, Leipzig, Germany

⁷ * Correspondence: [Sebastian Sosa <s.sosa@live.fr>](mailto:S Sebastian Sosa <s.sosa@live.fr>)

⁸ wordcount : 2959

⁹

Abstract

- ¹⁰ 1. Bayesian modeling is a powerful paradigm in modern statistics and machine learning, offering a principled framework for inference under uncertainty. However, practitioners face significant obstacles, including **interoperability** issues, a persistent **accessibility-flexibility trade-off**, the limitations of **domain-specific limitations**, and challenges in **scalability**.
- ¹⁵ 2. **Interoperability:** The landscape of Bayesian software is fragmented across programming languages and abstraction levels. Newcomers often gravitate towards high-level interfaces (e.g., *brms*) within familiar environments due to their accessibility for standard models. However, highgh levels of abstraction frameworks can be restrictive, lacking the flexibility needed for custom or complex models as research needs evolve.
- ²⁰ 3. **Accessibility-Flexibility Trade-off:** To gain the necessary flexibility, researchers must often transition to lower-level probabilistic programming languages . This transition imposes a steeper learning curve and requires mastery of specific modeling languages or complex programming frameworks, hindering broader adoption and rapid iteration.

- 25 4. ***Domain-Specific Limitations:*** Similar accessibility and flexibility trade-offs ex-
26 ist in domain-specific Bayesian packages. While providing accessible, pre-packaged
27 models for specific fields, customizing or extending these models often requires deep
28 engagement with lower-level programming languages or switching tools entirely, lim-
29 iting methodological innovation within those domains.
- 30 5. ***Scalability:*** Computational demands remain a significant bottleneck, limiting the
31 application of Bayesian methods to the large datasets and complex, high-dimensional
32 models prevalent in modern research.
- 33 6. To address these challenges, we introduce ***Bayesian Inference (BI)***, a new
34 Bayesian modeling software available in both Python and R. It aims to unify the
35 modeling experience by integrating an intuitive model-building syntax (enhancing
36 **accessibility**) with the **flexibility** of low-level abstraction coding available but also
37 pre-build function for high-level of abstraction and including hardware-accelerated
38 computation via JAX for improved **scalability**. Its availability in both major data
39 science languages directly tackles the **interoperability** barrier and the prebuild
40 function for specialized model in network analysis, survival models and phylogenetic
41 analysis allow to improved **domain-specific limitations**.

42

Introduction

43 Bayesian modeling has emerged as a vital tool in modern statistics and machine learning,
44 providing a framework for robust inference under uncertainty and the possibility to inte-
45 grate prior knowledge. Despite its potential, the practical application of Bayesian methods
46 is often hindered by significant hurdles within the current software ecosystem, preventing re-
47 searchers from fully leveraging its capabilities. Key challenges stem from the fragmented na-
48 ture of software across different programming languages (**interoperability**), gaps between
49 theoretical understanding and practical implementation (**accessibility**), complexities in
50 model specification that force trade-offs between ease-of-use and flexibility (**accessibility-**
51 **flexibility trade-off**), the constraints of overly specialized tools (**domain-specific limi-**
52 **tations**), and persistent computational scalability limitations for complex models or large
53 datasets (**scalability**).

54 The first major obstacle is the fragmented landscape of Bayesian software, scattered across
55 different programming languages and varying levels of abstraction, posing significant **in-**
56 **teroperability** challenges. Researchers frequently encounter a disparate collection of
57 tools—from *Stan*'s domain-specific language (DSL) to distinct low-level of abstraction
58 libraries (like *PyMC* (Salvatier, Wiecki, and Fonnesbeck 2016), *TensorFlow Probability*
59 (*TFP*) (Abadi et al. 2015), *NumPyro* (Phan, Pradhan, and Jankowiak 2019)) and high-
60 level of abstraction libraries (like *BRMS*). This fragmentation complicates workflows and

61 presents a confusing landscape, especially for researchers new to Bayesian analysis. For
62 instance, researchers new to Bayesian analysis may initially gravitate towards tools of
63 high-level of abstraction available within their most familiar programming environment
64 (e.g., *BRMS* (Bürkner 2017) in *R* (Wickham 2015)), potentially overlooking more suitable
65 options elsewhere due to the steep initial learning curve or perceived incompatibility (ac-
66 cessibility). This linguistic and platform diversity imposes considerable cognitive overhead,
67 potentially hindering the adoption of the most suitable tool for a given problem due to
68 familiarity biases or the friction of switching ecosystems, ultimately impacting the effec-
69 tive application of Bayesian methods. This initial hurdle of navigating disparate systems
70 naturally leads new practitioners to prioritize tools that appear easiest to learn, raising
71 concerns about the balance between accessibility and the flexibility needed for complex
72 research.

73 Compounding this fragmentation is the challenge of accessibility and the translation of
74 theoretical knowledge into practice **accessibility-flexibility trade-off**. Indeed, while
75 high-level interfaces like *BRMS* offer an intuitive formula-based syntax, significantly low-
76 ering the initial barrier to entry (e.g. generalized linear mixed models using *BRMS*), this
77 accessibility often comes at the cost of flexibility. As research questions become more
78 sophisticated, requiring custom likelihood functions (e.g., multiple likelihoods), intricate
79 prior structures (e.g., XXX), or non-standard model components (e.g., centered-random
80 factors), the limitations of these high-level wrappers become apparent. To gain the neces-
81 sary expressive power, the researcher must typically transition to lower-level probabilistic
82 programming languages (PPLs) such as *Stan* (Stan Development Team) (requiring mas-
83 tery of its specific DSL), *PyMC*, *NumPyro*, or *TFP*. This transition imposes a much steeper
84 learning curve, demanding a deeper understanding of probabilistic programming concepts
85 (like computational graphs or tensor manipulation) and often more verbose code. This
86 significant jump in complexity can deter users, divert focus from statistical modeling to
87 software engineering challenges, and ultimately slow down the pace of research, particularly
88 when trying to adapt models within specific scientific fields.

89 Similar accessibility and flexibility constraints manifest as **domain-specific limitations**
90 within specialized Bayesian packages. Fields like phylogenetics or network analysis benefit
91 from tools such as *BEAST* (Bouckaert 2019), *RevBayes* (Höhna et al. 2016), *STRAND*
92 (Ross, McElreath, and Redhead 2024), or *BISON* (Hart et al. 2023), which provide acces-
93 sible, pre-packaged models tailored to common domain problems. A phylogeneticist might
94 initially find *BEAST* convenient for standard molecular clock models. However, when
95 they wish to incorporate a novel evolutionary hypothesis requiring modification of the core
96 model structure or integrate data types not originally envisioned by the developers, they
97 often encounter rigid constraints. Extending these specialized tools frequently requires
98 deep engagement with their underlying, often complex, codebase (sometimes necessitat-
99 ing proficiency in languages like Java or C++) or abandoning the domain-specific tool
100 entirely in favor of a general-purpose PPL. This forces researchers to either compromise

101 on their methodological innovation or undertake a significant software development effort,
102 potentially switching programming ecosystems and losing the initial convenience, thereby
103 limiting the evolution of modeling practices within specialized domains. Even when model
104 specification is achievable, either in general or specialized tools, the computational feasi-
105 bility remains a major concern.

106 Finally, computational **scalability** continues to be a significant bottleneck, limiting the
107 application of Bayesian methods to the large datasets (e.g., millions of observations) and
108 complex, high-dimensional models (e.g., thousands of parameters) prevalent in modern re-
109 search across fields like genomics, neuroscience, and machine learning. While established
110 tools like *Stan* feature highly optimized inference algorithms (particularly its NUTS sam-
111 pler) and offer effective multi-core parallelization, they can still face challenges with long
112 C++ compilation times for complex models and may require substantial code restructur-
113 ing or external tooling to efficiently leverage hardware accelerators like GPUs or TPUs for
114 certain computations. Conversely, emerging frameworks built on *JAX* (Bradbury et al.
115 2018) (powering *NumPyro* and parts of *TFP*) promise substantial speedups via automatic
116 differentiation, JIT compilation, and native support for parallel hardware architectures.
117 However, integrating these powerful backends seamlessly into user-friendly, flexible model-
118 ing front-ends that don't require deep expertise in the *JAX* ecosystem itself is an ongoing
119 challenge. Domain-specific tools often inherit the scalability limitations of the frameworks
120 they are built upon, failing to provide a universally efficient solution across different model
121 types and data sizes.

122 Therefore, there is an evident and pressing demand for a Bayesian modeling framework
123 that synergistically addresses these interconnected limitations. To address these intercon-
124 nected challenges, we introduce ***Bayesian Inference (BI***, a new Bayesian modeling
125 software designed to unify the modeling experience across the two dominant data science
126 languages, Python and R. *BI* tackles the **interoperability** barrier head-on by offering
127 native interfaces in both environments. It aims to resolve the **accessibility-flexibility**
128 **trade-off** by providing an intuitive model-building syntax familiar to users of statistical
129 modeling languages, while enabling advanced customization and leveraging multiple, inter-
130 changeable inference backends for flexibility. To combat **domain-specific limitations**,
131 *BI* includes pre-built functions and structures tailored for specialized models in areas like
132 network analysis, survival analysis, and phylogenetic analysis, while still allowing exten-
133 sion and modification within its general framework. Crucially, *BI* enhances **scalability**
134 by integrating with hardware-accelerated computation via *JAX* (using *NumPyro* or *TFP*
135 as backends), enabling efficient execution on CPUs, GPUs, and TPUs. By providing a
136 streamlined, efficient, and unified environment for the end-to-end Bayesian workflow—
137 from model specification and fitting to diagnostics and prediction—***BI*** lowers the barrier
138 to entry for sophisticated Bayesian modeling, aiming to empower a broader community
139 of researchers across disciplines to confidently apply advanced Bayesian methods to their
140 complex research problems.

141 **Software Presentation**

142 *BI* directly confronts the **interoperability** challenge by offering native, feature-equivalent
143 implementations in both Python and R. While minor syntactic differences exist to ad-
144 here to the idiomatic conventions of each language, the core model specification syntax,
145 the procedural workflow for analysis, and the underlying computational engines remain
146 fundamentally consistent. For instance, Python utilizes dot notation for method calls on
147 class objects (e.g., `bi.dist.normal(0,1)`), while R employs dollar sign notation for ac-
148 cessing elements or methods within its object system (e.g., `bi$dist$normal(0,1)`). This
149 dual-language availability significantly lowers the adoption barrier for researchers, allowing
150 them to work entirely within their preferred programming environment without sacrificing
151 access to a common, powerful Bayesian modeling framework.

152 *BI* is designed to navigate the critical *accessibility-flexibility trade-off* by providing multiple
153 layers of abstraction and utility, catering effectively to users with varying levels of Bayesian
154 modeling expertise and diverse complexity requirements through : simplified backend inter-
155 action via intuitive syntax, pre-built components for complex model features , addressing
156 domain-specific limitations within a general framework, integrated End-to-End Workflow
157 and extensive model library and documentation.

158 At its computational core, *BI* leverages the power and efficiency of established Probabilistic
159 Programming Languages (PPLs) like *NumPyro* and *TFP*, both of which are built upon the
160 *JAX* framework for high-performance numerical computation and automatic differentiation.
161 However, *BI* deliberately abstracts away much of the inherent complexity of these lower-
162 level tools (**Code block 1**). This significantly enhances **accessibility** for a broader range
163 of users.

164

165 **Code block 1:** *Prior specification differences between NumPyro, TFP, and BI*

```
166 # NumPyro prior specification
167 numpyro.sample("mu", dist.Normal(0, 1)).expand([10])

# TFP prior specification (within a JointDistributionCoroutine)
yield Root(tfd.Sample(tfd.Normal(loc=1.0, scale=1.0), sample_shape=10))

# BI prior specification
bi.dist.normal(0, 1, name = "mu", shape = (10,))
```

168 To enhance **flexibility** without unduly sacrificing the accessibility provided by the high-
169 level syntax, *BI* includes a library of pre-built, computationally optimized functions imple-
170 mented directly in *JAX* (e.g., **Code block 2**). These components encapsulate common but

171 potentially complex modeling structures, allowing users to incorporate them easily within
172 the model specification. Key examples include:

- 173 1. *Centered Random Effects* and *Non-Centered Random Effects* for hierarchical (multi-
174 level) model components (McElreath 2018). The non-centered parameterization, of-
175 ten crucial for efficient sampling in hierarchical models (particularly with sparse data),
176 is provided without requiring the user to manually implement the reparameterization
177 logic.
- 178 2. *Kernels for Gaussian Processes* for modeling spatial, temporal, phylogenetic, or other
179 forms of structured correlation or dependency.
- 180 3. *Block Model Effects* for implementing stochastic block models in network analysis.
- 181 4. *SRM effects* for modeling pairwise interactions in networks while accounting for
182 sender effects, receiver effects, dyadic effects, nodal predictors, dyadic predictors,
183 and observation biases (Sosa et al., n.d.).
- 184 5. *Network-Based Diffusion Approach (NBDA)* components for modeling the effect of
185 network edges on the rates of transmission of phenomena (e.g., behavioral, epidemi-
186 ological) while accounting for nodal or dyadic covariates.
- 187 6. *Network metrics* ranging from nodal, dyadic, and global network measures with a
188 total of 11 that can be used to build custom models of social network analysis (Sosa,
189 Sueur, and Puga-Gonzalez 2020).

190 These pre-built *JAX* functions provide tailored model components for common patterns in
191 specific fields, while keeping them fully integrated within the general, extensible modeling
192 framework. By providing these optimized building blocks within its general syntax, *BI*
193 allows researchers in these fields to rapidly implement standard domain models using fa-
194 miliar concepts. Crucially, however, users retain the full flexibility of the *BI* framework to
195 combine these domain-specific components with other model features (e.g., complex non-
196 linear effects via splines, hierarchical structures across groups of networks or phylogenies)
197 or to customize or extend them using *BI*'s underlying mechanisms if needed—a capabili-
198 ty often missing in more narrowly focused domain-specific packages. This design aims
199 to foster methodological innovation *within* specialized domains by lowering the barrier to
200 implementing more complex or novel models [link to latex block].

201
202 **Code block 2:** Random effect specification differences between NumPyro, TFP, and BI
203

```
# NumPyro version of random centered effect
a = numpyro.sample("a", dist.Normal(5, 2))
    b = numpyro.sample("b", dist.Normal(-1, 0.5))
    sigma_cafe = numpyro.sample("sigma_cafe", dist.Exponential(1).expand([2]))
```

```

sigma = numpyro.sample("sigma", dist.Exponential(1))
Rho = numpyro.sample("Rho", dist.LKJ(2, 2))
cov = jnp.outer(sigma_cafe, sigma_cafe) * Rho
a_cafe_b_cafe = numpyro.sample(
    "a_cafe,b_cafe",
    dist.MultivariateNormal(jnp.stack([a, b]), cov).expand([20])
)
a_cafe, b_cafe = a_cafe_b_cafe[:, 0], a_cafe_b_cafe[:, 1]

# TFP version of random centered effect
alpha = yield Root(tfd.Sample(tfd.Normal(loc=5.0, scale=2.0), sample_shape=1))
beta = yield Root(tfd.Sample(tfd.Normal(loc=-1.0, scale=0.5), sample_shape=1))
sigma = yield Root(tfd.Sample(tfd.Exponential(rate=1.0), sample_shape=1))
sigma_alpha_beta = yield Root(tfd.Sample(tfd.Exponential(rate=1.0),
sample_shape=2))
Rho = yield Root(tfd.LKJ(dimension=2, concentration=2.0))
Mu = tf.concat([alpha, beta], axis=-1)
scale = tf.linalg.LinearOperatorDiag(sigma_alpha_beta).matmul(tf.squeeze(Rho))

# BI version of random centered effect
Sigma = dist.exponential(1, (ni,), name = 'Sigma_individual')
L = dist.lkjcholesky(1, (ni,), name = 'L_individual', shape = (ni,))
Z = dist.normal(0, 1, name = 'z_individual', shape = (ni,K))
alpha = random_centered2(Sigma, L, Z)

```

205 *BI* is designed to encapsulate the entire Bayesian modeling workflow within a cohesive
 206 object-oriented structure, promoting a streamlined and reproducible analysis pipeline. Typically,
 207 a user interacts with a primary *BI* object, through which they can sequentially:

- 208 • **Handle Data:** Load, preprocess, and associate dataset(s) with the model object.
- 209 • **Define Model:** Specify the model structure, including the likelihood(s), priors for
 210 all parameters, and incorporate any pre-built components using an intuitive formula
 syntax.
- 212 • **Run Inference:** Execute the model fitting process using the No-U-Turn Sampler
 213 (NUTS), which triggers the backend PPL (e.g., *NumPyro*, *TFP*) to perform Markov
 214 Chain Monte Carlo (MCMC) sampling. Progress indicators and diagnostics are typ-
 ically provided.

- **Analyze Posterior:** Access, summarize, and diagnose the posterior distributions of parameters. This includes methods for calculating posterior means, medians, credible intervals, convergence diagnostics (e.g., \hat{R} , Effective Sample Size - ESS), and retrieving raw posterior samples for custom analysis.
- **Visualize Results:** Generate standard diagnostic plots (e.g., trace plots, rank plots, posterior distributions) and visualizations of model parameters, effects, and predictions using integrated plotting functions that leverage the *arviz* library.

This unified structure minimizes the need for users to juggle multiple disparate software tools or manually transfer data and results between different stages of the analysis, thereby enhancing efficiency and reproducibility.

Finally, *BI* includes over 21 well-documented implementations of various standard and advanced Bayesian models. Examples include Generalized Linear Models (GLMs), Generalized Linear Mixed Models (GLMMs), survival analysis models (e.g., Cox proportional hazards), Principal Component Analysis (PCA), phylogenetic comparative methods, and various network models. Each implementation is accompanied by detailed documentation that encompasses: 1) general principles, 2) underlying assumptions, 3) code snippets in Python and R, and 4) mathematical details, enabling users to gain a deeper understanding of the modeling process and its nuances. Additionally, the framework's flexibility allows models to be combined; for example, building a zero-inflated model with varying intercepts and slopes, or constructing a joint model where principal components (derived from PCA) serve as predictors in a subsequent regression, allowing uncertainty to be propagated through all stages of the analysis.

Example : SRM model

To illustrate how these design features of *BI* coalesce to provide a streamlined, flexible, and powerful solution, effectively addressing the limitations identified in the existing Bayesian software landscape we will provide a basic example of how an SRM model is declared in *BI*, compare it with the equivalent model in Numpyro (Appendix 1) and STAN (Appendix 2). We will also show how this model can be build from scratch with *BI* (**Code block 3**) or its custom functions (**Code block 4**) to highlight the accessibility-flexibility of our package by demonstrating how advance user can build custom model (with less code than STAN) as well as how new user can apply pre-build *BI* models. Finally we show how it is also called in R (**Code block 5**) to cross language use with *BI*. Readers interested in further details on data structure, data import, data manipulation, and model fitting for SRM models can refer directly to the *BI* documentation [Modeling Network](#).

Code block 3: SRM model from scratch with BI

```

252
253 def model(N_id, idx, result_outcomes,
           focal_individual_predictors,
           target_individual_predictors):

    # Intercept
    intercept = bi.dist.normal(
        logit(0.1/jnp.sqrt(N_id)),
        2.5, shape=(1,), name = 'intercept'
    )

    # Sender receiver -----
    N_var = focal_individual_predictors.shape[0]
    N_id = focal_individual_predictors.shape[1]
    focal_effects = dist.normal(0, 1, name = 'focal_effects')
    target_effects = dist.normal( 0, 1, name = 'target_effects')
    terms = jnp.stack([
        focal_effects @ focal_individual_predictors,
        target_effects @ target_individual_predictors
    ], axis = -1)
    sr_raw = dist.normal(0, 1, shape=(2, N_id), name = 'sr_raw')
    sr_sigma = dist.exponential( 1, shape= (2,), name = 'sr_sigma')
    sr_L = dist.lkjcholesky(2, 2, name = "sr_L")
    rf = deterministic('sr_rf',(((sr_L @ sr_raw).T * sr_sigma)))
    ids = jnp.arange(0,sr_effects.shape[0])
    edgl_idx = bi.net.vec_node_to_edgle(jnp.stack([ids, ids], axis = -1))
    sender = sr_effects[edgl_idx[:,0],0] + sr_effects[edgl_idx[:,1],1]
    receiver = sr_effects[edgl_idx[:,1],0] + sr_effects[edgl_idx[:,0],1]
    sr = jnp.stack([sender, receiver], axis = 1)

    # dyadic effects -----
    bi.net.mat_to_edgl(dyadic_effect_mat)
    dr_raw = dist.normal(0, 1, shape=(2,N_dyads), name = 'dr_raw')
    dr_sigma = dist.exponential(1, name = 'dr_sigma' )
    dr_L = dist.lkjcholesky(2, 2, name = 'dr_L')
    dr_rf = deterministic('dr_rf', (
        ((dr_L @ dr_raw).T * jnp.repeat(dr_sigma, 2))
    ))

    dyad_effects = dist.normal(0, 1,
                               name= 'dyad_effects', shape = (dyadic_predictors.ndim - 1,

```

```

    ))
dr = dyad_effects * dyadic_predictors

# Likelihood
bi.dist.poisson(jnp.exp(intercept + sr + dr), obs=result_outcomes)

```

254

Code block 4: SRM model with prebuild functions

```

255 def model(N_id, idx, result_outcomes,
256         focal_individual_predictors,
257         target_individual_predictors):
258
259     # Intercept
260     intercept = bi.dist.normal(
261         logit(0.1/jnp.sqrt(N_id)),
262         2.5, shape=(1,), name = 'intercept'
263     )
264
265     # SR
266     sr = bi.net.sender_receiver(
267         focal_individual_predictors,
268         target_individual_predictors
269     )
270
271     # Dyadic
272     dr = bi.net.dyadic_effect(shape = idx.shape[0])
273
274     # Likelihood
275     bi.dist.poisson(jnp.exp(intercept + sr + dr), obs=result_outcomes)

```

258

Code block 5: SRM model with prebuild functions

```

259 model <- function(N_id, idxShape, result_outcomes,
260                     focal_individual_predictors, target_individual_predictors){
261
262     x=0.1/jnp$sqrt(N_id)
263     tmp=jnp$log(x / (1 - x))
264
265     # Intercept

```

```

intercept = bi.dist.normal(tmp, 2.5, shape=c(1), name = 'block')

# SR
sr = m$net$sender_receiver(
    focal_individual_predictors,
    target_individual_predictors
)

# Dyadic
dr = m$net$dyadic_effect(shape = c(idxShape))

# Likelihood
m$poisson(jnp$exp(intercept + sr + dr), obs=result_outcomes)
}

```

262 Finally, regarding code performance we can time the computation time for network of size
 263 200 in STAN and BI and observed that BI comput time is around XXX on cpu and XXX
 264 on gpu and STAN compute time around XXX.

265 **Discussion**

266 **BI** framework is built on top of the popular Python programming language, with a fo-
 267 cus on providing a user-friendly interface for model development and interpretation. Our
 268 framework is designed to be modular and extensible, allowing users to easily incorporate
 269 their own custom models and data types into the framework. One of the key features of
 270 this software is its comprehensive library of 21 predefined Bayesian models, covering a wide
 271 range of common applications and use cases. These models are accompanied by detailed
 272 explanations, making it easier for users to understand the underlying assumptions and
 273 apply the models to their specific research questions. In addition to these built-in models,
 274 the software includes several custom functions tailored for advanced statistical and network
 275 modeling. This curated library serves not only as a collection of ready-to-use tools but also
 276 as a valuable pedagogical resource, demonstrating best practices for constructing, fitting,
 277 and interpreting models within the *BI* framework, and providing robust templates for users
 278 aiming to develop novel model variants. Whether users are interested in hierarchical mod-
 279 els, time-series analysis, or cutting-edge network modeling approaches, our library caters
 280 to a variety of analytical needs. This accessibility fosters an environment where users can
 281 confidently explore and implement Bayesian methods, ultimately enhancing their research
 282 capabilities.

283 By providing a streamlined and efficient environment for the end-to-end Bayesian workflow—
284 from model specification and fitting to diagnostics and prediction, *BI* lowers the barrier
285 to entry for sophisticated Bayesian modeling. We aim to empower a broader community
286 of researchers across disciplines to confidently apply advanced Bayesian methods to their
287 complex research problems.

288 **References**

- 289 Abadi, Martín, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro,
290 Greg S. Corrado, et al. 2015. “TensorFlow: Large-Scale Machine Learning on Hetero-
291 geneous Systems.” <https://www.tensorflow.org/>.
- 292 Bouckaert, Timothy G. AND Barido-Sottani, Remco AND Vaughan. 2019. “BEAST 2.5:
293 An Advanced Software Platform for Bayesian Evolutionary Analysis.” *PLOS Compu-
294 tational Biology* 15 (4): 1–28. <https://doi.org/10.1371/journal.pcbi.1006650>.
- 295 Bradbury, James, Roy Frostig, Peter Hawkins, Matthew James Johnson, Chris Leary,
296 Dougal Maclaurin, George Necula, et al. 2018. “JAX: Composable Transformations of
297 Python+NumPy Programs.” <http://github.com/jax-ml/jax>.
- 298 Bürkner, Paul-Christian. 2017. “brms: An R Package for Bayesian Multilevel Models
299 Using Stan.” *Journal of Statistical Software* 80 (1): 1–28. [https://doi.org/10.18637/jss.v080.i01](https://doi.org/10.18637/jss.
300 v080.i01).
- 301 Hart, Jordan, Michael Nash Weiss, Daniel Franks, and Lauren Brent. 2023. “BISoN:
302 A Bayesian Framework for Inference of Social Networks.” *Methods in Ecology and
303 Evolution* 14 (9): 2411–20. <https://doi.org/10.1111/2041-210X.14171>.
- 304 Höhna, Sebastian, Michael J. Landis, Tracy A. Heath, Bastien Boussau, Nicolas Lartillot,
305 Brian R. Moore, John P. Huelsenbeck, and Fredrik Ronquist. 2016. “RevBayes:
306 Bayesian Phylogenetic Inference Using Graphical Models and an Interactive Model-
307 Specification Language.” *Systematic Biology* 65 (4): 726–36. [https://doi.org/10.1093/sysbio/syw021](https://doi.org/10.1093/
308 sysbio/syw021).
- 309 McElreath, Richard. 2018. *Statistical Rethinking: A Bayesian Course with Examples in r
310 and Stan*. Chapman; Hall/CRC.
- 311 Phan, Du, Neeraj Pradhan, and Martin Jankowiak. 2019. “Composable Effects for
312 Flexible and Accelerated Probabilistic Programming in NumPyro.” *arXiv Preprint
313 arXiv:1912.11554*.
- 314 Ross, Cody T, Richard McElreath, and Daniel Redhead. 2024. “Modelling Animal Network
315 Data in r Using STRAND.” *Journal of Animal Ecology* 93 (3): 254–66.
- 316 Salvatier, John, Thomas V Wiecki, and Christopher Fonnesbeck. 2016. “Probabilistic
317 Programming in Python Using PyMC3.” *PeerJ Computer Science* 2: e55.
- 318 Sosa, Sebastian, Mary B. McElreath, Daniel Redhead, and Cody T. Ross. n.d. “Robust
319 Bayesian Analysis of Animal Networks Subject to Biases in Sampling Intensity and

- 320 Censoring.” *Methods in Ecology and Evolution* n/a (n/a). <https://doi.org/https://doi.org/10.1111/2041-210X.70017>.
- 321 Sosa, Sebastian, Cédric Sueur, and Ivan Puga-Gonzalez. 2020. “Network Measures in
- 322 Animal Social Network Analysis: Their Strengths, Limits, Interpretations and Uses.”
- 323 Stan Development Team. “Stan Modeling Language Users Guide and Reference Manual,
- 324 Version 2.32.” <https://mc-stan.org>.
- 325 Wickham, Hadley. 2015. *R Packages*. 1st ed. O'Reilly Media, Inc.